

Build Task — Config & Logging Foundation — Hints & Solution

[Back to the Build Task](#) · [Hint 1](#) · [Hint 2](#) · [Solution](#)

Where this Build Task is saved: its own folder, `practice/build_task/` — the reusable foundation every later document copies:

```
practice/build_task/
├─ exceptions.py      MissingConfigError
├─ config.py          load_config() -> Config
├─ logging_setup.py   get_logger(name)
├─ test_config.py     proves the 4 Test Cases from the README
├─ .env.example       key names only, no real values (committed)
└─ .env              your real values (never committed)
```

Run it: `cd practice/build_task && python test_config.py` — from inside the folder, so `from config import load_config` finds the file next to it.

Builds on: `practice/config_logging_wiring/` ([Real-world exercise](#)) — **copy** that whole folder to `practice/build_task/`, then upgrade it: typed `Config`, a `.env.example`, a duplicate-handler guard in `get_logger()`, and `test_config.py` in place of `main.py`.

Used later by: every later document imports these two files — `practice/build_task/config.py` (`load_config()`) and `practice/build_task/logging_setup.py` (`get_logger(name)`). [Doc02](#) is the first to say so out loud, and copies both into its own project folder. The names `load_config()`, `Config`, `MissingConfigError` and `get_logger(name)` are fixed for exactly that reason.

Only 2 hints — work through them in order, and don't jump ahead until you've genuinely tried. Each hint has 2 depth levels: **Basic** (the plain idea) and **Intermediate** (proper Python — this is the level the Build Task expects). Read Basic first even if you already know Python — it's the fastest way to spot exactly what Intermediate adds.

- [Hint 1 — What you're building, and the exact pieces](#)
- [Hint 2 — The plan, and almost the whole thing](#)
- [Solution](#)

[Back to the Build Task](#) · [Hint 1](#) · [Hint 2](#) · [Solution](#)

Hint 1 — What you're building, and the exact pieces

Basic Version

You're building two small files that do the real work: `practice/build_task/config.py`, which loads settings, and `practice/build_task/logging_setup.py`, which sets up logging. Two more keep them honest: a one-line `practice/build_task/exceptions.py`, and `practice/build_task/test_config.py`, which proves the whole thing works.

The config loader's job: read some values from the environment, and either hand back something you can trust, or stop with a clear error message if something needed is missing. Fail early — right at startup — not later when it's confusing.

The logger's job: one function, `get_logger(name)`, that any file can call to get a working logger. Python's built-in `logging` module already does most of the work.

Here are the exact pieces you need to look up and use:

- `os.getenv(key)` — gives you the value, or `None` if it isn't set.
- `load_dotenv()` from `python-dotenv` — call it once near the top, before reading any values. It loads your `.env` file's contents in.
- A simple class (or a `@dataclass`) to hold your config values together.
- `logging.getLogger(name)` — always gives back the *same* logger if you call it twice with the same name.
- `logging.StreamHandler()` — prints to the screen. `logging.FileHandler(path)` — writes to a file.

[Back to the Build Task](#) · [Hint 1](#) · [Hint 2](#) · [Solution](#)

Intermediate Version

You're building two small, separate things that get used together: a **config loader** (`practice/build_task/config.py`) and a **logger factory** (`practice/build_task/logging_setup.py`). Keep them as two separate files, doing two separate jobs.

Think about the config loader first. Its whole job is: read some environment variables, and either hand back an object you can trust, or fail loudly with a clear reason. The key idea is **fail at startup, not later** — if `OPENAI_API_KEY` is missing, you want to know *immediately*, with a message that names the missing key, not five minutes later as a confusing `AuthenticationError` from OpenAI.

For the logger, the idea is: one function, `get_logger(name)` , that any file can call to get a working logger. Don't overthink this one — Python's built-in `logging` module already does almost everything you need.

Here's what to actually go look at:

- `os.getenv(key)` — returns the value, or `None` if the key isn't set. This is your basic tool for reading environment variables. `os.environ[key]` is the alternative, but it raises `KeyError` instead of returning `None` — decide which behavior you want to build on top of.
- `python-dotenv` 's `load_dotenv()` — call this once, near the top of your program, before you read any environment variables. It reads `.env` and loads its contents into `os.environ` for you.
- A **dataclass** (`from dataclasses import dataclass`) is a clean way to build your `Config` object — it gives you a typed object with almost no boilerplate.
- `logging.getLogger(name)` always returns the *same* logger object if called twice with the same name — Python caches them for you. You don't need to build your own caching.
- `logging.StreamHandler()` writes to the console; `logging.FileHandler(path)` writes to a file. Each handler can have its own `.setLevel(...)` , independent of the logger's own level.

Difference between Basic and Intermediate: Basic names the two files, their jobs, and the exact tools you'll need — `os.getenv` , `load_dotenv()` , a dataclass, `logging.getLogger` , and the two handler types — in plain words. Intermediate explains what each of those tools actually does and why, plus the "fail at startup" principle with a concrete reason it matters — this is the depth the Build Task's Solution is written at.

[Back to the Build Task](#) · [Hint 1](#) · [Hint 2](#) · [Solution](#)

Hint 2 — The plan, and almost the whole thing

Basic Version

```
practice/build_task/exceptions.py:
    make a MissingConfigError, it's a kind of Exception

practice/build_task/config.py:
    make a Config holder with fields: openai_api_key, log_level

    function load_config():
        run load_dotenv() to read the .env file
        read OPENAI_API_KEY
        if it's missing or empty: raise MissingConfigError, name the missing key
        read LOG_LEVEL, use "INFO" if it isn't set
        return a Config built from these

practice/build_task/logging_setup.py:
    function get_logger(name):
        get (or make) a logger for this name
        if it doesn't have a screen handler yet:
            make one, set its level from an env var (default INFO)
            attach it
        return the logger
```

The trickiest part of turning that plan into real code — checking a required key is really missing, not just empty:

```
# practice/build_task/config.py
import os

def require_env(key):
    value = os.getenv(key)
    if value is None or value == "":
        raise MissingConfigError("Required environment variable is missing: " + key)
    return value
```

Expected output if you run just this (nothing calls `require_env` yet): nothing — defining a function doesn't run it. You need to add a call below this to see anything happen.

Use this same idea for every required key. For the logger's double-handler problem:

```
# practice/build_task/logging_setup.py
if not logger.handlers:
    # only set up a handler the first time
    ...
```

Try finishing the rest yourself before looking at the Solution.

[Back to the Build Task](#) · [Hint 1](#) · [Hint 2](#) · [Solution](#)

Intermediate Version

practice/build_task/exceptions.py :

```
class MissingConfigError is an Exception
```

practice/build_task/config.py :

```
define a Config data holder with fields: openai_api_key, log_level

function load_config():
    call load_dotenv() to read the .env file into the environment
    read OPENAI_API_KEY from the environment
    if it is missing or empty:
        raise MissingConfigError, message naming "OPENAI_API_KEY"
    read LOG_LEVEL from the environment, default to "INFO" if not set
    return a Config built from these values
```

practice/build_task/logging_setup.py :

```
function get_logger(name):
    get (or create) a logger for this name
    if this logger doesn't already have a console handler attached:
        create a StreamHandler, set its level from an env var (default INFO)
        attach it to the logger
    set the logger's overall level low enough that the handler's level is what ac
    return the logger
```

The "if this logger doesn't already have a handler" check matters — without it, calling `get_logger("x")` twice would attach two handlers, and you'd see every log line printed twice.

Turning that plan into real code, here's the one piece worth seeing on its own first — this is *not* the full solution, just the trickiest part (checking a required key is genuinely missing, not just falsy):

```
# practice/build_task/config.py
import os

def require_env(key: str) -> str:
    value = os.getenv(key)
    if value is None or value == "":
```

```
    raise MissingConfigError(f"Required environment variable is missing: {key}")
    return value
```

Use this same pattern for every required key in `load_config()`, instead of writing the same `if` check out by hand each time.

For the logger's duplicate-handler problem, the check looks like this:

```
# practice/build_task/logging_setup.py
if not logger.handlers:
    # only attach a handler the first time this logger is configured
    ...
```

Try finishing the rest yourself before looking at the full Solution below.

Difference between Basic and Intermediate: Basic's pseudocode does the work every single time it's called, with no reuse, and its near-complete code proves the required-key check works at all. Intermediate is the same pseudocode shape described more precisely, plus the type contract real Python expects (`key: str -> str`) — this is the depth the Solution below is written at.

[Back to the Build Task](#) · [Hint 1](#) · [Hint 2](#) · [Solution](#)

Solution

Every code block below shows the exact output you'd see if you ran it, right after the code, on a machine with a `.env` file containing `OPENAI_API_KEY=sk-test-123` (and no `LOG_LEVEL`, unless a block says otherwise). Read both depths — they're not "wrong, right," they're 2 real, valid ways to solve the same problem, with real tradeoffs between them. This Build Task deliberately stops at Intermediate — it's the foundation every later document and project copies, so it stays the simplest version that's genuinely correct, not the fanciest one possible.

Everything below lives in one folder, and every code block names its file on the first line:

```
practice/build_task/
├─ exceptions.py      MissingConfigError
├─ config.py          load_config() -> Config
├─ logging_setup.py   get_logger(name)
├─ test_config.py     proves the Test Cases – run this one
├─ .env.example       key names only, no real values (committed)
└─ .env               your real values (never committed)
```

Files appear in the order you create them: `.env.example`, `.env`, `exceptions.py`, `config.py`, `logging_setup.py`, `test_config.py`. Run everything from inside the folder — `cd practice/build_task && python test_config.py` — so `from config import load_config` resolves.

Basic Version

Approach 1 — the direct way

```
# practice/build_task/.env.example
# Copy this file to .env and fill in the real values. Never commit .env.
OPENAI_API_KEY=
LOG_LEVEL=INFO
```

```
# practice/build_task/.env
OPENAI_API_KEY=sk-test-123
```

```
# practice/build_task/exceptions.py
class MissingConfigError(Exception):
    pass
```

```
# practice/build_task/config.py
import os
from dotenv import load_dotenv
from exceptions import MissingConfigError

class Config:
    def __init__(self, openai_api_key, log_level):
        self.openai_api_key = openai_api_key
        self.log_level = log_level

def require_env(key):
    value = os.getenv(key)
    if value is None or value == "":
        raise MissingConfigError("Required environment variable is missing: " + key)
    return value

def load_config():
    load_dotenv()
    api_key = require_env("OPENAI_API_KEY")
    log_level = os.getenv("LOG_LEVEL", "INFO")
    return Config(api_key, log_level)
```

```
# practice/build_task/logging_setup.py
import logging
import os

def get_logger(name):
    logger = logging.getLogger(name)
    if not logger.handlers:
        level_name = os.getenv("LOG_LEVEL", "INFO")
        level = getattr(logging, level_name.upper(), logging.INFO)
        handler = logging.StreamHandler()
        handler.setLevel(level)
        logger.addHandler(handler)
        logger.setLevel(level)
    return logger
```

```
# practice/build_task/test_config.py — proving it works
from config import load_config
from logging_setup import get_logger
```

```
config = load_config()
logger = get_logger(__name__)
logger.info("Loaded config for " + config.openai_api_key)
```

Expected output:

```
Loaded config for sk-test-123
```

And with `OPENAI_API_KEY` missing from `.env` entirely:

```
Traceback (most recent call last):
...
exceptions.MissingConfigError: Required environment variable is missing: OPENAI_A
```

This version works correctly and meets every Build Task requirement. It's missing type hints, and it builds error messages with `+` instead of an f-string — both fine for a first working version.

[Back to the Build Task](#) · [Hint 1](#) · [Hint 2](#) · [Solution](#)

Intermediate Version

Approach 1 — a dataclass-based config

`practice/build_task/.env.example` and `practice/build_task/.env` are exactly as shown in Basic Approach 1 above — keep both files as they are.

```
# practice/build_task/exceptions.py
class MissingConfigError(Exception):
    """Raised when a required setting is missing from the environment."""
    pass
```

```
# practice/build_task/config.py
import os
from dataclasses import dataclass
from dotenv import load_dotenv
from exceptions import MissingConfigError

@dataclass
class Config:
    openai_api_key: str
    log_level: str

def require_env(key: str) -> str:
    value = os.getenv(key)
    if value is None or value == "":
        raise MissingConfigError(f"Required environment variable is missing: {key}")
    return value

def load_config() -> Config:
    load_dotenv()
    api_key = require_env("OPENAI_API_KEY")
    log_level = os.getenv("LOG_LEVEL", "INFO")
    return Config(openai_api_key=api_key, log_level=log_level)
```

```
# practice/build_task/logging_setup.py
import logging
```

```

import os

def get_logger(name: str) -> logging.Logger:
    logger = logging.getLogger(name)

    if not logger.handlers:
        level_name = os.getenv("LOG_LEVEL", "INFO")
        level = getattr(logging, level_name.upper(), logging.INFO)

        handler = logging.StreamHandler()
        handler.setLevel(level)
        formatter = logging.Formatter("%(asctime)s %(name)s %(levelname)s %(message)s")
        handler.setFormatter(formatter)

        logger.addHandler(handler)
        logger.setLevel(level)

    return logger

```

```

# practice/build_task/test_config.py
"""Runs the four Test Cases from the README against config.py and logging_setup.py

Run it from inside this folder: python test_config.py
It rewrites .env as it goes, so it backs your real one up first and restores it afterwards
"""

import os
from pathlib import Path

from config import load_config
from exceptions import MissingConfigError
from logging_setup import get_logger

ENV_PATH = Path(".env")
BACKUP_PATH = Path(".env.backup")

def set_env_file(contents: str | None) -> None:
    """Put .env into a known state for one case, and clear os.environ first.

    load_dotenv() never overwrites a variable that is already set, so without
    these pop() calls every case after the first would inherit the one before it.
    """

    os.environ.pop("OPENAI_API_KEY", None)
    os.environ.pop("LOG_LEVEL", None)
    if contents is None:

```

```

        ENV_PATH.unlink(missing_ok=True)
    else:
        ENV_PATH.write_text(contents)

def main() -> None:
    if ENV_PATH.exists():
        ENV_PATH.rename(BACKUP_PATH)
    try:
        set_env_file(None)
        try:
            load_config()
        except MissingConfigError as e:
            print(f"case 1 (.env missing)      -> MissingConfigError: {e}")

        set_env_file("LOG_LEVEL=INFO\n")
        try:
            load_config()
        except MissingConfigError as e:
            print(f"case 2 (key missing)      -> MissingConfigError: {e}")

        set_env_file("OPENAI_API_KEY=sk-test-123\n")
        config = load_config()
        print(f"case 3 (valid .env)          -> Config loaded, log level {config.log

        get_logger("x")
        logger = get_logger("x")
        logger.info("case 4 (get_logger twice) -> printed exactly once")
    finally:
        ENV_PATH.unlink(missing_ok=True)
        if BACKUP_PATH.exists():
            BACKUP_PATH.rename(ENV_PATH)

if __name__ == "__main__":
    main()

```

Expected output (the exact timestamp will differ on your machine — `%(asctime)s` always prints the current time):

```

case 1 (.env missing)      -> MissingConfigError: Required environment variable is
case 2 (key missing)      -> MissingConfigError: Required environment variable is
case 3 (valid .env)       -> Config loaded, log level INFO
2026-09-10 09:00:00,123 x INFO case 4 (get_logger twice) -> printed exactly once

```

The last line is the only one that goes through the logger, so it's the only one carrying the timestamp/name/level prefix — and it appears once, not twice, which is the whole point of the `if not logger.handlers:` guard. One thing that surprises people: a `StreamHandler` writes to `stderr`, while `print()` writes to `stdout`. In a terminal they interleave in the order shown; pipe the output to a file (`python test_config.py > out.txt`) and the logger line can jump ahead of the others, because the two streams are buffered differently. Nothing is wrong when that happens.

Why this approach: a `@dataclass` gives you a typed, readable `Config` object with almost no extra code. This is a very common, standard pattern in real Python projects.

Approach 2 — a plain class, with file logging added

This version uses a plain class instead of a dataclass (more explicit, a bit more typing), and adds a file handler alongside the console one, since some projects want both.

`practice/build_task/exceptions.py` , `.env.example` and `.env` are unchanged from Approach 1 above — keep those three files exactly as they are. Only `config.py` , `logging_setup.py` and `test_config.py` change.

```
# practice/build_task/config.py
import os
from dotenv import load_dotenv
from exceptions import MissingConfigError

class Config:
    def __init__(self, openai_api_key: str, log_level: str) -> None:
        self.openai_api_key = openai_api_key
        self.log_level = log_level

def load_config() -> Config:
    load_dotenv()

    api_key = os.getenv("OPENAI_API_KEY")
    if not api_key:
        raise MissingConfigError("Required environment variable is missing: OPENAI_API_KEY")

    log_level = os.getenv("LOG_LEVEL", "INFO")
    return Config(openai_api_key=api_key, log_level=log_level)
```

```
# practice/build_task/logging_setup.py
import logging
import os
```

```
def get_logger(name: str) -> logging.Logger:
    logger = logging.getLogger(name)

    if logger.handlers:
        return logger

    console_level_name = os.getenv("LOG_LEVEL", "INFO")
    console_level = getattr(logging, console_level_name.upper(), logging.INFO)

    console_handler = logging.StreamHandler()
    console_handler.setLevel(console_level)

    file_handler = logging.FileHandler("app.log")
    file_handler.setLevel(logging.DEBUG)

    formatter = logging.Formatter("%(asctime)s %(name)s %(levelname)s %(message)s")
    console_handler.setFormatter(formatter)
    file_handler.setFormatter(formatter)

    logger.addHandler(console_handler)
    logger.addHandler(file_handler)
    logger.setLevel(logging.DEBUG)

    return logger
```

```
# practice/build_task/test_config.py
from config import load_config
from logging_setup import get_logger

config = load_config()
logger = get_logger(__name__)
logger.debug("Starting up")
logger.info(f"Loaded config for {config.openai_api_key}")
```

Expected output on screen (console handler is at `INFO` , so the `DEBUG` line is filtered out and never reaches the screen):

```
2026-09-10 09:00:00,456 __main__ INFO Loaded config for sk-test-123
```

Expected content appended to `app.log` (file handler is at `DEBUG` , so both lines land there):

```
2026-09-10 09:00:00,321 __main__ DEBUG Starting up
2026-09-10 09:00:00,456 __main__ INFO Loaded config for sk-test-123
```

Difference from Basic: both Intermediate approaches add full type hints — `openai_api_key: str`, `-> Config`, `-> logging.Logger` — turning every signature into documentation a reader (or an editor's autocomplete) can rely on without opening the function body. Both also add a real `Formatter`, so log lines carry a timestamp, logger name, and level instead of just the bare message. Approach 2 additionally adds a `FileHandler`, so you keep a full `DEBUG`-level record in `app.log` even while the screen only shows `INFO` and above — closer to what a real deployed script needs, since you can't always watch the screen live.

Which one should you actually use? Approach 1 (dataclass, console logging only) is what most people should ship — it satisfies every Build Task requirement with the least code. Approach 2's `FileHandler` is worth adding the moment you can't always watch the screen live for this program. Either is genuinely correct: this Build Task is the reusable foundation every later document and project copies, so it deliberately stops here instead of growing into a caching layer or a validation library — those are real, useful upgrades, but they're a choice to make later, in whichever specific project actually needs one, not a default every copy of this foundation should carry.