

Failure (respecting a real rate limit) — Solution

[Back to the exercise](#) · [Hint 1](#) · [Hint 2](#) · [Solution](#)

Basic Version

Approach 1 — the direct way

```
class FakeResponse:
    def __init__(self, status_code, headers):
        self.status_code = status_code
        self.headers = headers

def decide_wait_seconds(response, attempt):
    retry_after = response.headers.get("Retry-After")
    if retry_after is not None:
        return int(retry_after)
    return 2 ** attempt

# test: server told us how long to wait
r1 = FakeResponse(429, {"Retry-After": "2"})
print(decide_wait_seconds(r1, attempt=0))    # 2

# test: server didn't say, fall back to backoff
r2 = FakeResponse(429, {})
print(decide_wait_seconds(r2, attempt=3))    # 8
```

Expected output:

```
2
8
```

This is correct and covers both cases. It's missing type hints and doesn't guard against a malformed `Retry-After` value (like a non-numeric string) — fine for a first pass, worth fixing once you know it's a real risk.

[Back to the exercise](#) · [Hint 1](#) · [Hint 2](#) · [Solution](#)

Intermediate Version

Approach 1 — a dataclass fake, and a safe conversion

```
from dataclasses import dataclass, field

@dataclass
class FakeResponse:
    status_code: int
    headers: dict[str, str] = field(default_factory=dict)

def decide_wait_seconds(response: FakeResponse, attempt: int) -> int:
    """Respect Retry-After if the server sent one; otherwise fall back to backoff
    retry_after = response.headers.get("Retry-After")
    if retry_after is not None:
        try:
            return int(retry_after)
        except ValueError:
            pass # malformed header – fall through to backoff instead of crashing
    return 2 ** attempt

if __name__ == "__main__":
    told_to_wait = FakeResponse(status_code=429, headers={"Retry-After": "2"})
    print(decide_wait_seconds(told_to_wait, attempt=0)) # 2

    no_header = FakeResponse(status_code=429, headers={})
    print(decide_wait_seconds(no_header, attempt=3)) # 8

    malformed_header = FakeResponse(status_code=429, headers={"Retry-After": "soo
    print(decide_wait_seconds(malformed_header, attempt=2)) # 4, falls back safe
```

Expected output:

```
2
8
4
```

Difference from Basic: using a `@dataclass` for `FakeResponse` gives you a typed, self-documenting test double with almost no code. Wrapping the `int(retry_after)` conversion in its own `try/except ValueError` handles a malformed or unexpected header value — a genuinely rate-limited API is exactly the situation where you don't want one more edge case crashing your retry logic instead of just falling back sensibly.

[Back to the exercise](#) · [Hint 1](#) · [Hint 2](#) · [Solution](#)

Advanced Version

Approach 1 — the HTTP-date form, and a hard cap

```
from dataclasses import dataclass, field
from datetime import datetime, timezone
from email.utils import parsedate_to_datetime

@dataclass
class FakeResponse:
    status_code: int
    headers: dict[str, str] = field(default_factory=dict)

MAX_RETRY_AFTER_SECONDS = 60.0

def decide_wait_seconds(response: FakeResponse, attempt: int) -> float:
    """Respect Retry-After (seconds or an HTTP-date), capped so a server can
    never make this client sleep an unreasonable amount of time."""
    retry_after = response.headers.get("Retry-After")
    if retry_after is None:
        return float(2 ** attempt)

    try:
        wait = float(retry_after)
    except ValueError:
        try:
            target_time = parsedate_to_datetime(retry_after)
            wait = max((target_time - datetime.now(timezone.utc)).total_seconds(), 0)
        except (TypeError, ValueError):
            wait = float(2 ** attempt)

    return min(wait, MAX_RETRY_AFTER_SECONDS)

if __name__ == "__main__":
    plain_seconds = FakeResponse(429, {"Retry-After": "2"})
    print(decide_wait_seconds(plain_seconds, attempt=0)) # 2.0
```

```

no_header = FakeResponse(429, {})
print(decide_wait_seconds(no_header, attempt=3)) # 8.0

garbage_header = FakeResponse(429, {"Retry-After": "sometime later"})
print(decide_wait_seconds(garbage_header, attempt=2)) # 4.0, falls back to b

absurd_header = FakeResponse(429, {"Retry-After": "999999999"})
print(decide_wait_seconds(absurd_header, attempt=0)) # 60.0, capped

```

Expected output:

```

2.0
8.0
4.0
60.0

```

The last line is the important one: a server claiming you should wait for over 31 years gets capped at 60 seconds instead of being honored literally.

Approach 2 — the same idea, factored so the cap and the parsing are independently testable

```

from datetime import datetime, timezone
from email.utils import parsedate_to_datetime

MAX_RETRY_AFTER_SECONDS = 60.0

def parse_retry_after(value: str) -> float | None:
    """Return seconds to wait from a Retry-After value, or None if it
    can't be parsed in either the numeric or HTTP-date form."""
    try:
        return float(value)
    except ValueError:
        pass

    try:
        target_time = parsedate_to_datetime(value)
        return max((target_time - datetime.now(timezone.utc)).total_seconds(), 0.)
    except (TypeError, ValueError):
        return None

def decide_wait_seconds(headers: dict[str, str], attempt: int) -> float:
    retry_after = headers.get("Retry-After")

```

```

wait = parse_retry_after(retry_after) if retry_after is not None else None
if wait is None:
    wait = float(2 ** attempt)
return min(wait, MAX_RETRY_AFTER_SECONDS)

# parse_retry_after can be tested directly, with plain strings, no fake response
print(parse_retry_after("2")) # 2.0
print(parse_retry_after("not a number or a date")) # None
print(decide_wait_seconds({"Retry-After": "999999999"}, 0)) # 60.0

```

Expected output:

```

2.0
None
60.0

```

Difference from Intermediate, and between these 2 Advanced approaches: Intermediate trusts `Retry-After` to always be a plain integer string from a well-behaved server — a date-formatted header or a deliberately huge number either crashes it or makes it wait an absurd amount of time. Approach 1 fixes both in one function. Approach 2 splits the parsing (`parse_retry_after`) from the capping-and-fallback decision (`decide_wait_seconds`), so the trickiest part — turning a raw header string into a number of seconds, in either of its two legal forms — can be tested directly with plain string inputs, the same "pull the risky logic into its own pure function" idea from the timeout/retry exercise's `compute_backoff_delay` .

Which one should you actually write? For `http_client.py` in this document's Build Task, Approach 2's split is worth it — `parse_retry_after()` is exactly the kind of function you want a handful of quick unit tests for (a plain number, an HTTP-date, garbage, `None`), and testing it doesn't require constructing a fake response object at all. The hard cap (`MAX_RETRY_AFTER_SECONDS`) is worth keeping in any version of this code that talks to a server you don't fully control — which, for an external API, is always.