

Edge cases (where the $\frac{3}{4}$ -word rule breaks) — Hints

[Back to the exercise](#) · [Hint 1](#) · [Hint 2](#) · [Solution](#)

Only 2 hints — work through them in order, and don't jump ahead until you've genuinely tried. Each hint has 2 depth levels: **Basic** (the plain idea) and **Intermediate** (why it happens). Read Basic first even if you already know the rule — it's the fastest way to spot exactly what Intermediate adds.

- [Hint 1 — Why the rule breaks, and where](#)
- [Hint 2 — Worked examples, and the real numbers](#)

[Back to the exercise](#) · [Hint 1](#) · [Hint 2](#) · [Solution](#)

Hint 1 — Why the rule breaks, and where

Basic Version

The "1 token $\approx \frac{3}{4}$ of a word" rule was worked out from ordinary English prose. It's not a law of nature — it's just a description of how the tokenizer happens to behave on the kind of text it saw a lot of during training. Anything that isn't ordinary English prose is fair game to break the rule badly.

Pick one snippet of Python code, and one sentence written in a language other than English, and just guess their token counts using the rule before you look anything up.

[Back to the exercise](#) · [Hint 1](#) · [Hint 2](#) · [Solution](#)

Intermediate Version

The tokenizer's vocabulary is built by a statistical process (byte-pair encoding, roughly: repeatedly merging the most frequent adjacent character pairs seen in training) run over a training corpus that's mostly English prose. Whatever pattern shows up often in that corpus earns a short, efficient token. Whatever doesn't — rare words, punctuation-heavy code syntax, non-Latin scripts — gets broken down into smaller, less efficient pieces, because the exact chunk wasn't common enough in training to deserve its own single token.

For code: count the punctuation characters separately from the words — parentheses, quotes, colons, and brackets each often become their own token, which is why code "looks short" but tokenizes long.

For non-English text: don't try to guess from an English rule of thumb at all here — just note that the same *meaning*, in a language with less training data behind it, usually costs noticeably more tokens than the English equivalent.

Write your guesses for the code snippet and the non-English sentence, in tokens, before checking them.

The deeper pattern connecting both exceptions (code, non-English text) is the same: **token efficiency tracks training-data frequency, not human-perceived complexity**. A short Urdu sentence isn't linguistically more complex than its English translation — it just has far less representation in the corpus the tokenizer's vocabulary was built from. Given this, predict: would a JSON-heavy system prompt (lots of `{`, `}`, `"`, `:`) tokenize better or worse than the equivalent information written as plain English sentences? Write your prediction, then check it against a real tokenizer.

Difference between Basic and Intermediate: Basic names that the rule breaks outside plain English. Intermediate explains the training-data mechanism behind why, and generalizes it into a rule you can apply to any new content type you haven't tested yet.

[Back to the exercise](#) · [Hint 1](#) · [Hint 2](#) · [Solution](#)

Hint 2 — Worked examples, and the real numbers

Basic Version

```
code snippet: "print('hello world')"  
word-count guess: about 3 tokens (2 words + quotes)  
actual: usually 5-7 tokens – each punctuation mark counts separately  
  
urdu sentence: "میں آج بازار جا رہا ہوں" (I am going to the market today)  
word-count guess: about 7 tokens (5 words)  
actual: usually much higher – often 2-3x the English equivalent
```

The point isn't to memorize a multiplier — it's to know when to stop trusting the quick word-count rule entirely and just measure.

[Back to the exercise](#) · [Hint 1](#) · [Hint 2](#) · [Solution](#)

Intermediate Version

```
code: "print('hello world')"
```

naive word-count estimate: ~3 tokens
 actual (tokenizer tool): ~5-7 tokens
 gap explained by: print, (, ', hello, world, ',) each frequently tokenizing separately or near-separately – punctuation isn't "free"


```
non-English: "میں آج بازار جا رہا ہوں"
```

English equivalent: "I am going to the market today" (~7 tokens)
 actual for the Urdu sentence: often 15-20+ tokens
 gap explained by: far less Urdu text in the tokenizer's training data means the vocabulary has few or no efficient multi-character chunks for this script – most of it falls back to small, inefficient sub-word or byte-level pieces

Run your own two examples through the real tokenizer tool now and record the actual counts. Write down which parts of your own future projects fall into these exception categories: system prompts full of JSON schema examples, a multilingual support bot, or a feature that echoes back user-pasted code, then compare against the [Solution](#).

Difference between Basic and Intermediate: Basic and Intermediate both get you accurate numbers for the 2 examples in front of you — Intermediate additionally explains the mechanism behind the gap, and has you generalize it to your own future projects.

[Back to the exercise](#) · [Hint 1](#) · [Hint 2](#) · [Solution](#)

Full solution: [Show me the solution](#)