

Edge cases (where the 3/4-word rule breaks) — Solution

[Back to the exercise](#) · [Hint 1](#) · [Hint 2](#) · [Solution](#)

Basic Version

Text	Word-count guess	Actual (tokenizer tool)
<code>print('hello world')</code>	~3 tokens	usually 5-7 tokens
"میں آج بازار جا رہا ہوں" (Urdu, "I am going to the market today")	~7 tokens	often 15-20+ tokens

Code costs more than it looks like because punctuation (`(` , `'` , `)`) usually tokenizes as its own piece, not "for free" the way a space between English words is. Non-English text costs more because the tokenizer's vocabulary was built mostly from English training data — a script with less representation there gets broken into smaller, less efficient pieces to represent the same meaning.

The one-line rule to actually use: if the content isn't plain English prose (code, another language, dense structured data like JSON), stop estimating and run it through the real tokenizer tool or `tiktoken`.

[Back to the exercise](#) · [Hint 1](#) · [Hint 2](#) · [Solution](#)

Intermediate Version

Code: `print('hello world')`

naive word-count estimate (2 "words" + punctuation) :	~3 tokens
actual	: 5-7 tokens

Likely token breakdown: `print` , `(` , `'` , `hello` , `world` , `'` , `)` — roughly one token per syntactic element, because code's punctuation carries meaning the tokenizer's training data treated as worth its own token, unlike a plain-English comma that usually merges into a neighboring word's token.

Non-English: "میں آج بازار جا رہا ہوں"

English-equivalent estimate ("I am going to the market today", ~7 tokens) :	~7 to
actual	: often

The mechanism: byte-pair-encoding-style tokenizers build their vocabulary from the most frequent character sequences in a training corpus. A corpus that's mostly English text produces a vocabulary rich in efficient whole-word and common-subword English tokens. A script underrepresented in that training data has few or no such efficient chunks reserved for it, so the tokenizer falls back to much smaller sub-word or even byte-level pieces to represent the same text.

Difference from Basic: Basic gives you the numbers and the gap. Intermediate explains the mechanism (byte-pair encoding trained on a mostly-English corpus) that produces that gap for both code and non-English text — the same underlying cause, showing up as two different-looking symptoms.

[Back to the exercise](#) · [Hint 1](#) · [Hint 2](#) · [Solution](#)

Advanced Version

Approach 1 — a reusable measurement helper

```
import tiktoken

def real_token_count(text: str, model: str = "gpt-4o") -> int:
    encoding = tiktoken.encoding_for_model(model)
    return len(encoding.encode(text))

print(real_token_count("print('hello world')"))      # actual count, not a guess
print(real_token_count("میں آج بازار جا رہا ہوں"))    # actual count, not a guess
```

This is the concrete answer to "stop estimating" — a one-line call any part of the codebase can use before shipping a system prompt, a few-shot example set, or a multilingual UI string, instead of relying on a developer remembering that this particular content type is an exception.

Approach 2 — measuring per-language cost from real traffic, not a guessed multiplier

```
def average_tokens_per_language(logged_messages: list[dict]) -> dict[str, float]:
    totals: dict[str, list[int]] = {}
    for msg in logged_messages:
        lang = msg["language"]
        totals.setdefault(lang, []).append(real_token_count(msg["text"]))
    return {lang: sum(counts) / len(counts) for lang, counts in totals.items()}
```

Why this matters beyond Approach 1: "non-English costs 2-3x more" is a useful rule of thumb for a quick gut-check, but real languages vary a lot — some non-Latin scripts cost far more than others, and the only way to know your actual per-language cost is to measure it from real traffic, not assume one multiplier covers every language your feature serves.

Why this matters practically, beyond trivia: a system prompt full of JSON-schema examples or embedded code tokenizes far worse than its word count suggests, inflating cost on every single call it's part of. A multilingual product priced the same per user regardless of language will quietly cost more per conversation for non-English users — worth knowing before you build a pricing model or a context-budget assumption around English-only testing.

Which one should you actually write? For exploring a new content type once, use the browser tokenizer tool — it's fast and needs no code. For anything shipped in a real feature, write Approach 1's `real_token_count()` helper and call it wherever a prompt, template, or user-facing string is built, so nobody has to remember the exception list by heart. If the feature serves multiple languages at real scale, add Approach 2's per-language measurement from actual logged traffic — a single "2-3x" rule of thumb is a fine estimate for day one, but real cost and budget decisions deserve real per-language numbers once you have production data to measure from.