

Intermediate (chunking method changes the answer) — Solution

[Back to the exercise](#) · [Hint 1](#) · [Hint 2](#) · [Solution](#)

Story — `chunking_practice.py`: the same document, split two different ways, can hand the model two different answers to the same question — this is the exercise that makes "chunking is a real design decision" something you watched happen, not something you took on faith. **If not:** the Build Task's chunker would be the first place a split fact ever bit you, in a pipeline with far more moving parts to debug through.

All examples below use the same test document — shown here as prose, not a code block, since the exact line-wrapping inside `DOCUMENT`'s triple-quoted string is what makes `chunk_by_chars` cut through "scalds" later on:

```
DOCUMENT (a 3-paragraph string, separated by blank lines) — "Coffee brewing starts with the grind size. A coarse grind works best for a French press, because its metal filter lets fine grounds slip through and makes the coffee taste muddy and bitter. A fine grind is right for espresso instead, where water touches the grounds for only a few seconds." / "Water temperature matters just as much as grind size. Water that is too hot, above 205 degrees Fahrenheit, scalds the coffee and pulls out bitter compounds that should have stayed in the grounds. Water that is too cool under-extracts the coffee, leaving it weak and sour even with plenty of grounds." / "The final step is the pour itself. A slow, steady pour in a spiral pattern wets all the grounds evenly, which is why baristas practice just the pour for weeks before worrying about anything else."
```

```
QUESTION = "What water temperature should be avoided when brewing coffee?"
```

Basic Version

Approach 1 — the direct way

```
# chunking_practice.py – Intermediate section
import math
from openai import OpenAI
```

```

client = OpenAI()

def chunk_by_chars(text, chunk_size=200):
    chunks = []
    for i in range(0, len(text), chunk_size):
        chunks.append(text[i:i+chunk_size])
    return chunks

def chunk_by_paragraph(text):
    chunks = []
    for piece in text.split("\n\n"):
        stripped = piece.strip()
        if stripped:
            chunks.append(stripped)
    return chunks

def get_embedding(text):
    response = client.embeddings.create(model="text-embedding-3-small", input=text)
    return response.data[0].embedding

def cosine_similarity(a, b):
    dot = sum(x * y for x, y in zip(a, b))
    norm_a = math.sqrt(sum(x * x for x in a))
    norm_b = math.sqrt(sum(x * x for x in b))
    return dot / (norm_a * norm_b)

def find_best_chunk(chunks, question_vector):
    best_chunk = None
    best_score = -1
    for chunk in chunks:
        chunk_vector = get_embedding(chunk)
        score = cosine_similarity(question_vector, chunk_vector)
        if score > best_score:
            best_score = score
            best_chunk = chunk
    return best_chunk, best_score

question_vector = get_embedding(QUESTION)

char_chunks = chunk_by_chars(DOCUMENT)
paragraph_chunks = chunk_by_paragraph(DOCUMENT)

char_best, char_score = find_best_chunk(char_chunks, question_vector)
paragraph_best, paragraph_score = find_best_chunk(paragraph_chunks, question_vector)

print("chunk_by_chars best match, score:", char_score)
print(char_best)

```

```
print()
print("chunk_by_paragraph best match, score:", paragraph_score)
print(paragraph_best)
```

Expected output: with `chunk_size=200`, `chunk_by_chars` cuts this document right through the word "scalds" — one chunk ends `"...205 degrees Fahrenheit, sc"` and the next one starts `"alds the coffee..."`. Whichever half wins, it's missing part of the answer.

`chunk_by_paragraph` keeps the whole "Water temperature matters..." paragraph — including "205 degrees Fahrenheit" and the reason why — as one unbroken chunk.

This version works correctly for what the exercise asks. It's missing type hints and reuses the same `get_embedding` call for every chunk on every run, with no reuse across calls.

[Back to the exercise](#) · [Hint 1](#) · [Hint 2](#) · [Solution](#)

Intermediate Version

Approach 1 — type hints and a `find_best_chunk` helper

```
# chunking_practice.py – Intermediate section
import math
from openai import OpenAI

client = OpenAI()

DOCUMENT: str = "..." # as above
QUESTION: str = "What water temperature should be avoided when brewing coffee?"

def chunk_by_chars(text: str, chunk_size: int = 200) -> list[str]:
    chunks = []
    for i in range(0, len(text), chunk_size):
        chunks.append(text[i:i + chunk_size])
    return chunks

def chunk_by_paragraph(text: str) -> list[str]:
    chunks = []
    for piece in text.split("\n\n"):
        stripped = piece.strip()
        if stripped:
            chunks.append(stripped)
    return chunks

def get_embedding(text: str) -> list[float]:
    response = client.embeddings.create(model="text-embedding-3-small", input=text)
    return response.data[0].embedding

def cosine_similarity(a: list[float], b: list[float]) -> float:
    dot = sum(x * y for x, y in zip(a, b))
    norm_a = math.sqrt(sum(x * x for x in a))
    norm_b = math.sqrt(sum(x * x for x in b))
    return dot / (norm_a * norm_b)
```

```

def find_best_chunk(
    chunks: list[str], question_vector: list[float],
) -> tuple[str, float]:
    best_chunk = ""
    best_score = -1.0
    for chunk in chunks:
        chunk_vector = get_embedding(chunk)
        score = cosine_similarity(question_vector, chunk_vector)
        if score > best_score:
            best_score = score
            best_chunk = chunk
    return best_chunk, best_score

def main() -> None:
    question_vector = get_embedding(QUESTION)
    methods = {
        "chunk_by_chars": chunk_by_chars(DOCUMENT),
        "chunk_by_paragraph": chunk_by_paragraph(DOCUMENT),
    }
    for name, chunks in methods.items():
        best_chunk, best_score = find_best_chunk(chunks, question_vector)
        print(f"{name} best match, score: {best_score:.4f}")
        print(best_chunk)
        print()

if __name__ == "__main__":
    main()

```

Expected output: same conclusion as Basic — `chunk_by_paragraph` 's top match contains the whole temperature fact; `chunk_by_chars` 's top match is missing half of "scalds."

Difference from Basic: full type hints on every function. `find_best_chunk` has an explicit `tuple[str, float]` return type, so any caller knows exactly what comes back without reading its body. The two methods are driven from one `methods` dict and one loop in `main()`, instead of copy-pasted print statements for each — adding a third chunking method later means adding one dict entry, not duplicating the comparison logic again.

Approach 2 — overlapping fixed-size chunks

Story: `chunk_by_chars` from Approach 1 can still cut a fact clean in half between two non-overlapping chunks — this is the actual fix, not a bigger `chunk_size` (which only moves the problem). **If not:** the `chunk_boundary_split` exercise's whole failure mode would still be sitting unfixed in this same file.

```
# chunking_practice.py – Intermediate section
def chunk_by_chars_with_overlap(
    text: str, chunk_size: int = 200, overlap: int = 40,
) -> list[str]:
    chunks: list[str] = []
    step = chunk_size - overlap # why: less than chunk_size, so chunks share text
    i = 0
    while i < len(text):
        chunks.append(text[i:i + chunk_size])
        i += step
    return chunks
```

Run this against `DOCUMENT` with `chunk_size=200`, `overlap=40` and look at the chunk boundaries: the word "scalds" that got split in half by the non-overlapping version now appears whole in at least one chunk, because each new chunk starts 40 characters *before* the previous one ended. The word that was previously exactly on the cut line is now safely inside the overlapping region instead of split by it.

Approach 3 — a paragraph chunker that doesn't silently give up

Story: `chunk_by_paragraph` trusts the document has real blank-line breaks — hand it one giant wall of text (a scraped page, a PDF that lost its formatting) and it silently returns the whole thing as "1 chunk," no error, no warning. **If not:** a search index that's secretly not chunked at all would look identical to a working one, until a search near the end of that one giant chunk came back oddly.

```
# chunking_practice.py – Intermediate section
def chunk_by_paragraph_safe(text: str) -> list[str]:
    pieces = chunk_by_paragraph(text)
    if len(pieces) <= 1:
        # how: no real paragraph breaks found -- fall back instead of
        # silently returning the whole document as "1 chunk"
        return chunk_by_chars_with_overlap(text, chunk_size=200, overlap=40)
    return pieces

# proof this matters: a document with no blank lines at all
single_block = DOCUMENT.replace("\n\n", " ")

naive_result = chunk_by_paragraph(single_block)
safe_result = chunk_by_paragraph_safe(single_block)

naive_count = len(naive_result)
safe_count = len(safe_result)
```

```
print("naive chunk_by_paragraph, no blank lines:", naive_count, "chunk(s)")
print("chunk_by_paragraph_safe, same document:", safe_count, "chunk(s)")
```

Expected output:

```
naive chunk_by_paragraph, no blank lines: 1 chunk(s)
chunk_by_paragraph_safe, same document: 4 chunk(s)
```

The plain `chunk_by_paragraph` hands back the entire document as a single "chunk" — no error, no warning, just a search index that's secretly not chunked at all. `chunk_by_paragraph_safe` notices that and falls back to the overlapping character chunker instead.

Difference from Approach 1, and between Approaches 2/3: Approach 1's two chunkers both produce clean, non-overlapping pieces, and `chunk_by_paragraph` trusts that the document has real paragraph breaks to split on. Approach 2 fixes the boundary problem directly: overlap means a fact sitting near a cut usually survives whole in at least one chunk, without needing a bigger `k` at search time to compensate (the workaround `chunk_boundary_split`'s solution describes). Approach 3 fixes a different, quieter failure — a document that defeats `chunk_by_paragraph` entirely by having no blank lines — by detecting that case and falling back to Approach 2's chunker instead of returning one useless giant chunk.

Which one should you actually write? For a document set you've actually looked at (like this exercise's `DOCUMENT`), Approach 1's plain `chunk_by_paragraph` is fine — real paragraph breaks, no boundary problem to speak of. The moment your documents come from somewhere you don't fully control — scraped pages, converted PDFs, user uploads — reach for Approach 3's fallback, since you can't guarantee every file will have the blank-line structure you're assuming. Reach for Approach 2's overlap specifically for `chunk_by_chars`-style fixed-size chunking, any time a fact could plausibly land near a chunk boundary — which, over a large enough document set, it eventually will.