

Real-world (build and search a real knowledge base) — Hints

[Back to the exercise](#) · [Hint 1](#) · [Hint 2](#) · [Solution](#)

Only 2 hints — work through them in order, and don't jump ahead until you've genuinely tried. Each hint has 3 depth levels: **Basic** (the plain idea), **Intermediate** (proper Python), **Advanced** (how a real knowledge base survives being rerun, and tells you when nothing good matched). Read Basic first even if you already know Python — it's the fastest way to spot exactly what each deeper level adds.

- [Hint 1 — The idea, and the exact pieces](#)
- [Hint 2 — The plan, and almost the whole thing](#)

[Back to the exercise](#) · [Hint 1](#) · [Hint 2](#) · [Solution](#)

Hint 1 — The idea, and the exact pieces

Basic Version

Pick a topic you actually know something about — coffee brewing, your favorite hobby, whatever. Write 5-10 short files about it, a few sentences each, and save them as plain text files in a folder.

Then turn each file into an embedding (a list of numbers that captures its meaning) and store those embeddings in a vector database. When you ask a question, you embed the question the same way, and ask the database which stored documents are the closest match.

The whole point of writing the documents yourself is that you already know the right answer. If you ask "what grind size should I use for a French press?" and the search doesn't return your French press doc, you know immediately something's wrong — you don't need to guess.

[Back to the exercise](#) · [Hint 1](#) · [Hint 2](#) · [Solution](#)

Intermediate Version

Same idea, naming the exact pieces:

- `client.embeddings.create(model="text-embedding-3-small", input=text)` — turns one file's text into a vector.
- `chromadb.Client()` — a local vector database, no server to set up.
- `chroma_client.create_collection(name="kb")` — a named table of documents plus their embeddings.
- `collection.add(documents=[...], embeddings=[...], ids=[...])` — three parallel lists: raw text, matching embedding, and a unique id (the filename works).
- `collection.query(query_embeddings=[question_embedding], n_results=3)` — note `query_embeddings`, not `query_texts`, since you embedded the question yourself.

Because you wrote every document, you're your own ground truth — you can check each search result by eye and know immediately whether it's right or wrong, no guessing required.

[Back to the exercise](#) · [Hint 1](#) · [Hint 2](#) · [Solution](#)

Advanced Version

Two things break the moment this stops being a script you run exactly once.

First: `chroma_client.create_collection(name="kb")` raises an error if a collection named "kb" already exists. Run your script a second time — to add a new document, or just because you re-ran it while testing — and it crashes on line 2, before doing anything useful. A script, a test file, and a notebook might all build "the kb collection" the same way this document's Build Task expects several files to share a config loader — and none of them should care whether it's the first call or the fiftieth. `chroma_client.get_or_create_collection(name="kb")` fixes this: it returns the existing collection if one's already there, and creates it only if it isn't — the same "check first, only do the real work if needed" idea from Doc01's Build Task cache.

Second: ask an off-topic question — something none of your documents are actually about — and `collection.query()` still confidently returns your `n_results` closest matches, because that's what a nearest-neighbor search always does: it finds the closest *available* things, even when none of them are actually close. Nothing in the code signals "none of these are good matches." A real system checks the returned distance/score against a cutoff, and either drops the weak results or tells the caller explicitly that nothing relevant was found — instead of quietly handing the model 3 irrelevant chunks and letting it guess from there.

The extra pieces:

- `get_or_create_collection(name="kb")` instead of `create_collection(name="kb")`.
- A `SCORE_CUTOFF` constant, and a filter step after `collection.query()` that drops any result whose distance is worse than the cutoff, before returning.

Sketch what a `search(question, cutoff)` function's signature and return type should look like once it can return "nothing relevant" as a real, distinct case — not just an empty list that looks the same as "no documents at all."

Difference between Basic, Intermediate, and Advanced: Basic and Intermediate both assume this script runs exactly once, against documents guaranteed to contain the answer to whatever you ask. Advanced removes both assumptions: `get_or_create_collection` makes building the store safe to call more than once, from more than one place, and a score cutoff makes "nothing relevant was found" a real, checked outcome instead of 3 confidently-wrong results that look no different from a genuine match.

[Back to the exercise](#) · [Hint 1](#) · [Hint 2](#) · [Solution](#)

Hint 2 — The plan, and almost the whole thing

Basic Version

```
pick a topic (coffee brewing at home)

write 6 short files, a few sentences each, save them to a docs folder:
    French press basics.txt
    Pour-over basics.txt
    Grind size guide.txt
    Water temperature.txt
    Coffee-to-water ratio.txt
    Storing coffee beans.txt

make a function called build_vector_store that takes a folder path:
    start a chroma client
    make a collection called "kb"
    for each file in the folder:
        read its full text
        embed the text
        add it to the collection, using the filename as the id
    return the collection

ask 3 real questions, one at a time:
    embed the question
    query the collection for the top 3 matches
    print the matched ids and text
    look at the results yourself and check: is this actually the right doc?
```

Here's almost the whole thing — try running it and reading it line by line:

```
import chromadb
from openai import OpenAI

openai_client = OpenAI()

def embed(text):
    response = openai_client.embeddings.create(model="text-embedding-3-small", in
    return response.data[0].embedding
```

```
def build_vector_store(doc_folder):  
    chroma_client = chromadb.Client()  
    collection = chroma_client.create_collection(name="kb")  
    # loop over the files in doc_folder here, embed each one, and  
    # collection.add(...) it using the filename as the id  
    return collection
```

Expected output if you run just this: nothing yet — write the 6 sample files, fill in the loop, and add the part that asks 3 questions and prints the results.

[Back to the exercise](#) · [Hint 1](#) · [Hint 2](#) · [Solution](#)

Intermediate Version

```

define:
    def build_vector_store(doc_folder: str):
        client = chromadb.Client()
        collection = client.create_collection(name="kb")
        for each file_path in Path(doc_folder).iterdir():
            text = file_path.read_text()
            embedding = embed(text)
            collection.add(documents=[text], embeddings=[embedding], ids=[file_pa
        return collection

define:
    def ask(collection, question: str) -> None:
        question_embedding = embed(question)
        results = collection.query(query_embeddings=[question_embedding], n_resul
        print the question, then results["ids"][0] and results["documents"][0]

test:
    collection = build_vector_store("docs")
    for question in your 3 questions:
        ask(collection, question)

```

Write `build_vector_store` 's loop and `ask` yourself, then compare against the [Solution](#).

[Back to the exercise](#) · [Hint 1](#) · [Hint 2](#) · [Solution](#)

Advanced Version

```
define:
    SCORE_CUTOFF = 0.35      # a distance worse than this means "not actually rele

    def build_vector_store(doc_folder):
        client = chromadb.Client()
        collection = client.get_or_create_collection(name="kb")    # safe to call
        if collection.count() == len(files in doc_folder): return collection    # a
        for each file: embed and add, same as before
        return collection

    def search(collection, question, k=3, cutoff=SCORE_CUTOFF):
        question_embedding = embed(question)
        results = collection.query(query_embeddings=[question_embedding], n_resul
        keep only results whose distance <= cutoff
        if nothing survives the cutoff: return [] and note "no relevant match"
        return the surviving (id, text, distance) tuples
```

Here's almost the whole thing — fill in the missing piece yourself:

```
SCORE_CUTOFF: float = 0.35

def build_vector_store(doc_folder: str):
    chroma_client = chromadb.Client()
    collection = chroma_client.get_or_create_collection(name="kb")
    # your turn: if this collection is already populated (collection.count() > 0)
    # just return it instead of re-embedding and re-adding every file again
    ...
    return collection

def search(collection, question: str, k: int = 3, cutoff: float = SCORE_CUTOFF) -
    question_embedding = embed(question)
    results = collection.query(query_embeddings=[question_embedding], n_results=k
    # your turn: build a list of {"id":..., "text":..., "score":...} dicts,
    # but only for results whose distance is <= cutoff -- drop the rest
    ...
```

Fill in the "already built" check and the cutoff filter yourself, then compare all 3 of your finished versions against the [Solution](#).

Difference between Basic, Intermediate, and Advanced: Basic and Intermediate build the store once and trust every query returns something worth using. Advanced makes `build_vector_store` safe to call again without duplicating or crashing, and makes `search` capable of reporting "nothing relevant" as a real, distinct result instead of silently returning 3 weak matches with no way to tell they're weak.

[Back to the exercise](#) · [Hint 1](#) · [Hint 2](#) · [Solution](#)

Full solution: [Show me the solution](#)