

Real-world (build and search a real knowledge base) — Solution

[Back to the exercise](#) · [Hint 1](#) · [Hint 2](#) · [Solution](#)

Every example below assumes the same 6 sample documents (`docs/*.txt`):

```
# knowledge_base_search_practice.py
DOCS: dict[str, str] = {
    "French press basics.txt": "A French press steeps coarse coffee grounds direc
    "Pour-over basics.txt": "Pour-over brewing means pouring hot water over groun
    "Grind size guide.txt": "Grind size should match your brew method. Use a coar
    "Water temperature.txt": "Water for brewing coffee should be between 195 and
    "Coffee-to-water ratio.txt": "A common starting ratio is 1 gram of coffee to
    "Storing coffee beans.txt": "Store coffee beans in an airtight container, awa
}
```

Basic Version

Approach 1 — raw `chromadb` client

```
# knowledge_base_search_practice.py
from pathlib import Path
import chromadb
from openai import OpenAI

openai_client = OpenAI()

docs_folder = Path("docs")
docs_folder.mkdir(exist_ok=True)
for filename, text in DOCS.items():
    (docs_folder / filename).write_text(text)

def embed(text):
    response = openai_client.embeddings.create(model="text-embedding-3-small", in
    return response.data[0].embedding

def build_vector_store(doc_folder):
```

```

chroma_client = chromadb.Client()
collection = chroma_client.create_collection(name="kb")
for file_path in Path(doc_folder).iterdir():
    text = file_path.read_text()
    embedding = embed(text)
    collection.add(documents=[text], embeddings=[embedding], ids=[file_path.name])
return collection

collection = build_vector_store("docs")

questions = [
    "What grind size should I use for a French press?",
    "How hot should the water be when brewing?",
    "How long does ground coffee stay fresh?",
]

for question in questions:
    question_embedding = embed(question)
    results = collection.query(query_embeddings=[question_embedding], n_results=3)
    print("Q: " + question)
    print(results["ids"][0])
    print(results["documents"][0])
    print()

```

Expected output: since you wrote every doc yourself, question 1 should return `French press basics.txt` (and probably `Grind size guide.txt`) at the top, question 2 should return `Water temperature.txt`, and question 3 should return `Storing coffee beans.txt`. Read the printed results and check.

This version works correctly and meets every requirement of the exercise. Running the script a second time crashes, though — `create_collection` raises if `"kb"` already exists.

[Back to the exercise](#) · [Hint 1](#) · [Hint 2](#) · [Solution](#)

Intermediate Version

Approach 1 — raw `chromadb`, factored into functions with type hints

```
# knowledge_base_search_practice.py
from pathlib import Path
import chromadb
from openai import OpenAI

openai_client = OpenAI()

def write_sample_docs(doc_folder: str, docs: dict[str, str]) -> None:
    folder_path = Path(doc_folder)
    folder_path.mkdir(exist_ok=True)
    for filename, text in docs.items():
        (folder_path / filename).write_text(text)

def embed(text: str) -> list[float]:
    response = openai_client.embeddings.create(model="text-embedding-3-small", input=text)
    return response.data[0].embedding

def build_vector_store(doc_folder: str):
    chroma_client = chromadb.Client()
    collection = chroma_client.create_collection(name="kb")
    for file_path in Path(doc_folder).iterdir():
        text = file_path.read_text()
        embedding = embed(text)
        collection.add(documents=[text], embeddings=[embedding], ids=[file_path.name])
    return collection

def ask(collection, question: str) -> None:
    question_embedding = embed(question)
    results = collection.query(query_embeddings=[question_embedding], n_results=3)
    print(f"Q: {question}")
    print(results["ids"][0])
    print(results["documents"][0])
    print()
```

```
def main() -> None:
    write_sample_docs("docs", DOCS)
    collection = build_vector_store("docs")
    questions = [
        "What grind size should I use for a French press?",
        "How hot should the water be when brewing?",
        "How long does ground coffee stay fresh?",
    ]
    for question in questions:
        ask(collection, question)

if __name__ == "__main__":
    main()
```

Approach 2 — LangChain's Chroma wrapper

```
# knowledge_base_search_practice.py
from langchain_chroma import Chroma
from langchain_openai import OpenAIEmbeddings

def build_vector_store(docs: dict[str, str]) -> Chroma:
    texts = list(docs.values())
    metadatas = [{"source": filename} for filename in docs.keys()]
    return Chroma.from_texts(texts, OpenAIEmbeddings(), metadatas=metadatas)

def ask(vector_store: Chroma, question: str) -> None:
    results = vector_store.similarity_search(question, k=3)
    print(f"Q: {question}")
    for result in results:
        print(result.metadata["source"], "- ", result.page_content)
    print()

def main() -> None:
    vector_store = build_vector_store(DOCS)
    questions = [
        "What grind size should I use for a French press?",
        "How hot should the water be when brewing?",
        "How long does ground coffee stay fresh?",
    ]
    for question in questions:
```

```
ask(vector_store, question)

if __name__ == "__main__":
    main()
```

Expected output: same 3 correct top matches as Basic. `Chroma.from_texts` embeds every doc for you, and `similarity_search` embeds the question for you — no manual `embeddings.create()` call anywhere.

Difference from Basic: full type hints, and everything split into named functions (`write_sample_docs`, `build_vector_store`, `ask`, `main`) instead of one flat script — each function does one job and can be tested or reused on its own. Approach 2 additionally removes the manual embedding calls entirely by letting LangChain's `Chroma` wrapper handle them. Neither approach yet survives being run a second time, or tells you when a question's best match genuinely isn't relevant — that's what Advanced adds.

[Back to the exercise](#) · [Hint 1](#) · [Hint 2](#) · [Solution](#)

Advanced Version

Approach 1 — `get_or_create_collection`, safe to call more than once

```
# knowledge_base_search_practice.py
from pathlib import Path
import chromadb
from openai import OpenAI

openai_client = OpenAI()

def embed(text: str) -> list[float]:
    response = openai_client.embeddings.create(model="text-embedding-3-small", in
    return response.data[0].embedding

def build_vector_store(doc_folder: str):
    chroma_client = chromadb.Client()
    collection = chroma_client.get_or_create_collection(name="kb")

    file_paths = list(Path(doc_folder).iterdir())
    if collection.count() >= len(file_paths):
        # already built by an earlier call in this same run -- don't re-embed eve
        return collection

    for file_path in file_paths:
        text = file_path.read_text()
        embedding = embed(text)
        collection.add(documents=[text], embeddings=[embedding], ids=[file_path.n
    return collection

# proof this is now safe to call twice
first = build_vector_store("docs")
second = build_vector_store("docs")
print(first.count(), second.count())
```

Expected output:

6 6

The second call doesn't crash, and doesn't double-add every document either — `collection.count() >= len(file_paths)` short-circuits it.

Approach 2 — a score cutoff, so "no relevant match" is a real, checked outcome

```
# knowledge_base_search_practice.py
SCORE_CUTOFF: float = 0.35 # tune this against your own embedding model + documents

def search(collection, question: str, k: int = 3, cutoff: float = SCORE_CUTOFF) -> list[dict]:
    question_embedding = embed(question)
    results = collection.query(query_embeddings=[question_embedding], n_results=k)

    matches: list[dict] = []
    for doc_id, text, distance in zip(
        results["ids"][0], results["documents"][0], results["distances"][0]
    ):
        if distance <= cutoff:
            matches.append({"id": doc_id, "text": text, "score": distance})

    return matches # empty list means "nothing relevant," not "search failed"

collection = build_vector_store("docs")

good_question = "What grind size should I use for a French press?"
off_topic_question = "What's the best way to train a marathon?"

print("on-topic:", search(collection, good_question))
print("off-topic:", search(collection, off_topic_question))
```

Expected output (exact scores vary, but the shape holds):

```
on-topic: [{'id': 'French press basics.txt', 'text': '...', 'score': 0.18}, ...]
off-topic: []
```

The on-topic question returns real matches under the cutoff. The off-topic question still gets a nearest neighbor internally — nearest-neighbor search always finds something — but every distance is worse than `cutoff`, so `search()` filters all of them out and returns `[]` instead of confidently handing back 3 irrelevant chunks.

Difference from Intermediate, and between these 2 Advanced approaches: Intermediate's `build_vector_store` crashes on a second call, and its `ask / similarity_search` always return *something*, whether or not it's actually relevant. Approach 1 fixes the first problem with `get_or_create_collection` plus a cheap "already built" check — the same "check first, only do the real work if needed" pattern as Doc01's `load_config()` cache. Approach 2 fixes a different problem: it makes "nothing here is actually relevant" a real, testable outcome instead of a silent 3-result list that looks identical whether the match was great or terrible.

Which one should you actually use? Both, together, are what a real knowledge base needs — they solve genuinely different problems. `get_or_create_collection` (Approach 1) costs nothing and should just always be there, the moment there's any chance this code runs more than once. The score cutoff (Approach 2) needs real tuning against your own documents and embedding model — a cutoff picked without ever checking a genuinely off-topic question against it is just a guess. Both directly feed the Build Task below: `retrieve()` 's "handles an empty set of documents without crashing" requirement is the same idea as `get_or_create_collection` 's safety, and "off-topic questions should return low scores... not confidently-wrong top results" is exactly what Approach 2's cutoff gives you.