

Failure (an endless loop, and a real pause/resume) — Hints

[Back to the exercise](#) · [Hint 1](#) · [Hint 2](#) · [Solution](#)

Only 2 hints, covering both halves of this exercise — the endless loop, and the checkpoint/interrupt/resume cycle. Work through them in order, and don't jump ahead until you've genuinely tried both halves. Each hint has 3 depth levels: **Basic** (the plain idea), **Intermediate** (proper LangGraph), **Advanced** (what actually gets persisted, and why that matters for a real pause). Read Basic first even if you already know LangGraph — it's the fastest way to spot exactly what each deeper level adds.

- [Hint 1 — The idea, and the exact pieces](#)
- [Hint 2 — The plan, and almost the whole thing](#)

[Back to the exercise](#) · [Hint 1](#) · [Hint 2](#) · [Solution](#)

Hint 1 — The idea, and the exact pieces

Basic Version

Part 1, the endless loop: build 2 nodes with a plain edge from each to the other (`a -> b` , `b -> a`) and no conditional edge that ever leaves the cycle. Run it with a low, hard step limit so it fails fast and clearly instead of actually hanging.

Part 2, checkpoint + interrupt + resume: a checkpointer is what makes a graph's state survive between two separate `invoke()` calls. `interrupt()` is a function you call inside a node that pauses the whole graph right there. You invoke once (it pauses), then invoke again later, on the same "thread," to pick up exactly where it stopped.

Things to use:

- `config={"recursion_limit": 5}` passed to `graph.invoke(...)` — the step-limit guard for Part 1.
- `from langgraph.checkpoint.memory import MemorySaver` and `builder.compile(checkpointer=MemorySaver())` — for Part 2.
- `from langgraph.types import interrupt, Command` — `interrupt(...)` pauses; `Command(resume=...)` is what you invoke with to continue.
- A `thread_id` inside `config={"configurable": {"thread_id": "..."}}` — this is what ties the two separate `invoke()` calls together as "the same paused run."

[Back to the exercise](#) · [Hint 1](#) · [Hint 2](#) · [Solution](#)

Intermediate Version

Part 1: `graph.invoke({...}, config={"recursion_limit": 5})` on a graph with no real exit raises `GraphRecursionError` once it hits the limit — this is the graph-level version of the step limit Doc07's `while` loop needed too. The point isn't to avoid ever seeing this error; it's to have actually watched it happen once, on purpose, so an accidental version of it in a real project isn't a surprise.

Part 2: the exact pieces:

- `checkpointer = MemorySaver()` , then `graph = builder.compile(checkpointer=checkpointer)` .
- A `thread_id` , e.g. `config = {"configurable": {"thread_id": "run-1"}}` — every `invoke()` using this same `config` reads and writes the *same* saved state.
- Inside one node: `from langgraph.types import interrupt` then `decision = interrupt({"question": "Approve this action?"})` — calling this pauses the graph right there and surfaces the payload you passed it to whoever's running the graph.
- First call: `result = graph.invoke({...}, config=config)` — this returns with the graph paused; `result` carries the interrupt payload (check `result["__interrupt__"]`), not a final answer yet.
- Resume call: `graph.invoke(Command(resume=True), config=config)` — same `config` , same `thread_id` , telling the paused `interrupt()` call what value to "return" so the node can keep going from exactly where it stopped.

[Back to the exercise](#) · [Hint 1](#) · [Hint 2](#) · [Solution](#)

Advanced Version

Think about what a checkpointer is actually persisting, and when. It's not just "the current value of one field" — it's the *entire* state, plus enough information for LangGraph to know exactly which node to run next. That's why a resumed run doesn't restart the node that called `interrupt()` from scratch with a blank state — it picks the *same* node back up, with everything that node's earlier code already did still intact, as if nothing had happened in between.

The real design question isn't just "does the state survive" — it's "would the *same* checkpointer setup actually survive your program restarting, not just a second function call in the same script?"

The extra piece that matters here:

- `MemorySaver()` only ever lives in your program's memory — it's genuinely gone the instant the Python process exits, exactly like Doc04's chatbot list. Proving pause/resume "works" by calling `invoke()` twice in the same script, with the same `MemorySaver` object still sitting in a variable, doesn't actually prove state survives a real restart — it only proves the object in memory still has it. A real production checkpointer (`langgraph.checkpoint.sqlite`, `langgraph.checkpoint.postgres`, etc.) writes to an actual database, so the *second* `invoke()` call could happen from an entirely separate program run, hours later, and still resume correctly.
- To genuinely test this like a restart, structure your resume as **a second, separate script run** (or at minimum, delete and recreate every Python variable except the `thread_id` between your "pause" and "resume" code) — if you resume using the exact same in-memory `checkpointer` object your first call already had a reference to, you haven't actually tested that the *state* persisted, only that the *object* did.

Difference between Basic, Intermediate, and Advanced: Basic and Intermediate get a working pause/resume cycle with `MemorySaver`, in one script. Advanced asks whether that test actually proves what it claims to — and points out that `MemorySaver`'s "memory" is exactly as fragile as Doc04's chatbot list unless you structure your resume test as a genuinely separate run, which is the only way to tell "the checkpointer object survived" apart from "the state actually survived."

[Back to the exercise](#) · [Hint 1](#) · [Hint 2](#) · [Solution](#)

Hint 2 — The plan, and almost the whole thing

Basic Version

Part 1 -- endless loop:

```
node_a: returns some small change
node_b: returns some small change
edges: start -> node_a -> node_b -> node_a (no exit, ever)
invoke with a low recursion_limit, inside try/except, and read the error
```

Part 2 -- checkpoint + interrupt + resume:

```
make a checkpointer (MemorySaver)
compile the graph with it
one node calls interrupt(...) partway through
pick a thread_id, put it in config
```

```
first invoke: graph pauses at interrupt() -- read what comes back
second invoke, same thread_id, using Command(resume=...): graph continues and
```

Here's almost the whole thing for Part 1 — just try running it and reading it line by line:

```
# loop_limit_interrupt_practice.py
from langgraph.graph import StateGraph, START, END
from langgraph.errors import GraphRecursionError

class LoopState(TypedDict):
    count: int

def node_a(state):
    return {"count": state["count"] + 1}

def node_b(state):
    return {"count": state["count"] + 1}

builder = StateGraph(LoopState)
builder.add_node("node_a", node_a)
builder.add_node("node_b", node_b)
builder.add_edge(START, "node_a")
```

```
builder.add_edge("node_a", "node_b")
builder.add_edge("node_b", "node_a") # no exit -- this is the endless loop, on p
```

Expected output if you run just this: nothing — add `compile()` and an `invoke({"count": 0}, config={"recursion_limit": 5})` inside a `try / except GraphRecursionError` to see it fail on purpose, fast.

[Back to the exercise](#) · [Hint 1](#) · [Hint 2](#) · [Solution](#)

Intermediate Version

Part 1, complete:

```
# loop_limit_interrupt_practice.py
from typing import TypedDict
from langgraph.graph import StateGraph, START, END
from langgraph.errors import GraphRecursionError

class LoopState(TypedDict):
    count: int

def node_a(state: LoopState) -> dict:
    return {"count": state["count"] + 1}

def node_b(state: LoopState) -> dict:
    return {"count": state["count"] + 1}

builder = StateGraph(LoopState)
builder.add_node("node_a", node_a)
builder.add_node("node_b", node_b)
builder.add_edge(START, "node_a")
builder.add_edge("node_a", "node_b")
builder.add_edge("node_b", "node_a")

graph = builder.compile()

try:
    graph.invoke({"count": 0}, config={"recursion_limit": 5})
except GraphRecursionError as e:
    print(f"Stopped on purpose: {e}")
```

Part 2, the plan for the rest:

```
from langgraph.checkpoint.memory import MemorySaver
from langgraph.types import interrupt, Command

class ApprovalState(TypedDict):
    task: str
    approved: bool
```

```

def do_task(state):
    decision = interrupt({"question": f"Approve running: {state['task']}?"})
    return {"approved": decision}

builder = StateGraph(ApprovalState)
builder.add_node("do_task", do_task)
builder.add_edge(START, "do_task")
builder.add_edge("do_task", END)

checkpointer = MemorySaver()
graph = builder.compile(checkpointer=checkpointer)

config = {"configurable": {"thread_id": "run-1"}}
first_result = graph.invoke({"task": "send an email", "approved": False}, config=
print(first_result) # paused -- look for the interrupt payload here

second_result = graph.invoke(Command(resume=True), config=config)
print(second_result) # resumed -- approved should now be True

```

Run this and check: does `first_result` actually contain the question you passed to `interrupt()` ? Does `second_result["approved"]` come back `True` ? Compare against the [Solution](#).

[Back to the exercise](#) · [Hint 1](#) · [Hint 2](#) · [Solution](#)

Advanced Version

same Part 2 graph and checkpointer setup, but split into 2 actual separate script sharing only the `thread_id` (written down or hardcoded), to prove the state really and isn't just living in a Python variable your script still happens to hold:

`pause_script.py`:

build the graph, compile with `MemorySaver`, invoke once with a fixed `thread_id`

`resume_script.py` (run separately, after `pause_script.py` has already exited):

build the SAME graph shape again (new `MemorySaver()` object -- but see the cat invoke with `Command(resume=...)` and the SAME `thread_id`, confirm it resumes co

Here's the catch worth noticing yourself before reading the Solution: with `MemorySaver`, a genuinely fresh `resume_script.py` process **can't** actually resume — because `MemorySaver`'s data really does live only in the process that created it, and a new process gets a brand-new, empty one. Try it and watch it fail (or simply not find the paused state) — that failure *is* the lesson: it's the proof that `MemorySaver` is exactly as fragile as Doc04's in-memory chatbot list, and that a checkpointer meant to survive a real restart needs to write to an actual file or database (`langgraph.checkpoint.sqlite.SqliteSaver`, for instance), not RAM.

Difference between Basic, Intermediate, and Advanced: Basic and Intermediate build a working pause/resume cycle and test it within one running script, using `MemorySaver`. Advanced deliberately tries to resume from a second, truly separate process, and treats the resulting failure as the real finding — proving hands-on that `MemorySaver` doesn't survive a restart, which is exactly why production systems need a persisted checkpointer instead.

[Back to the exercise](#) · [Hint 1](#) · [Hint 2](#) · [Solution](#)

Full solution: [Show me the solution](#)