

Real-world (build and compare the supervisor version) — Solution

[Back to the exercise](#) · [Hint 1](#) · [Hint 2](#) · [Solution](#)

Story — `architecture_comparison_practice.py` (Real-world section): the supervisor version of the same task, measured the same way, so the "why a supervisor?" question gets a real answer instead of a feeling. **If not:** Project 4's "Plan Before You Code" step would defend its design with guesses.

Both depths below reuse `model`, `research()`, `write()`, and `StageResult` from `sequential_measure`'s solution unchanged — they sit higher up in this same file. Read both depths — they're not "wrong, right," they're 2 real, valid ways to solve the same problem, with real tradeoffs between them.

Basic Version

Approach 1 — the direct way, `if`-based routing

Story: the smallest working supervisor — plain `if` checks deciding the next node — built on the exact same `research()` and `write()` as the sequential version, so only the routing differs. **If not:** a difference in the numbers could come from different stage code, not from the design.

```
# architecture_comparison_practice.py – Real-world section
from typing import TypedDict
from langgraph.graph import StateGraph, START, END
from langgraph.types import Command

class SupervisorState(TypedDict):
    task: str
    research_text: str
    summary: str
    routing_log: list
    total_tokens: int
    total_seconds: float

def supervisor(state):
```

```

# how: plain checks on state – "which stage has no output yet?"
if not state.get("research_text"):
    log = state["routing_log"] + ["research"]
    return Command(update={"routing_log": log}, goto="research_node")
if not state.get("summary"):
    log = state["routing_log"] + ["write"]
    return Command(update={"routing_log": log}, goto="write_node")
return Command(goto=END)

def research_node(state):
    result = research(state["task"])
    return Command(
        update={
            "research_text": result.text,
            "total_tokens": state["total_tokens"] + result.tokens,
            "total_seconds": state["total_seconds"] + result.seconds,
        },
        goto="supervisor",
    )

def write_node(state):
    result = write(state["research_text"])
    return Command(
        update={
            "summary": result.text,
            "total_tokens": state["total_tokens"] + result.tokens,
            "total_seconds": state["total_seconds"] + result.seconds,
        },
        goto="supervisor",
    )

builder = StateGraph(SupervisorState)
builder.add_node("supervisor", supervisor)
builder.add_node("research_node", research_node)
builder.add_node("write_node", write_node)
builder.add_edge(START, "supervisor")
graph = builder.compile()

# when: only when this file is run directly – so another practice file
# can import these functions without triggering a paid run
if __name__ == "__main__":
    result = graph.invoke({
        "task": "climate change",
        "research_text": "",
    })

```

```
        "summary": "",
        "routing_log": [],
        "total_tokens": 0,
        "total_seconds": 0.0,
    })
    print(result["routing_log"])
    print(result["total_tokens"], "tokens,", result["total_seconds"],
          "seconds")
```

Expected output (exact numbers vary by run):

```
['research', 'write']
306 tokens, 2.81 seconds
```

This works and correctly reuses `sequential_measure`'s functions unchanged. It's missing a fair, run-both comparison against the sequential pipeline, and its routing decision costs nothing — both fine for a first working version.

[Back to the exercise](#) · [Hint 1](#) · [Hint 2](#) · [Solution](#)

Intermediate Version

Approach 1 — a `run_supervisor()` report matching `run_sequential()`'s shape

Story: a comparison is only fair when both sides report the same fields; `run_supervisor()` mirrors `run_sequential()` exactly. **If not:** you'd be comparing numbers that were collected in different ways.

```
# architecture_comparison_practice.py — Real-world section
def run_supervisor(topic: str) -> dict:
    initial_state = {
        "task": topic,
        "research_text": "",
        "summary": "",
        "routing_log": [],
        "total_tokens": 0,
        "total_seconds": 0.0,
    }
    final_state = graph.invoke(initial_state)
    return {
        "routing_log": final_state["routing_log"],
        "total_tokens": final_state["total_tokens"],
        "total_seconds": final_state["total_seconds"],
    }

if __name__ == "__main__":
    topics = ["climate change", "inflation", "sleep quality"]
    for topic in topics:
        sequential_report = run_sequential(topic)    # from sequential_measure
        supervisor_report = run_supervisor(topic)    # this exercise

        seq_tokens = sequential_report["total_tokens"]
        seq_seconds = sequential_report["total_seconds"]
        sup_tokens = supervisor_report["total_tokens"]
        sup_seconds = supervisor_report["total_seconds"]
        sup_log = supervisor_report["routing_log"]

        print(f"\n{topic}")
        print(f"  sequential: {seq_tokens} tokens, {seq_seconds:.2f}s")
```

```
print(f"  supervisor: {sup_tokens} tokens, {sup_seconds:.2f}s")
print(f"  routed: {sup_log}")
```

Expected output (exact numbers vary by run):

```
climate change
  sequential: 306 tokens, 2.75s
  supervisor: 306 tokens, 2.81s
  routed: ['research', 'write']

inflation
  sequential: 286 tokens, 2.60s
  supervisor: 286 tokens, 2.66s
  routed: ['research', 'write']

sleep quality
  sequential: 325 tokens, 2.90s
  supervisor: 325 tokens, 2.96s
  routed: ['research', 'write']
```

Difference from Basic: `run_supervisor()` returns the exact same 3 keys (`total_tokens`, `total_seconds`, plus `routing_log` in place of nothing) as `sequential_measure`'s `run_sequential()`, so the two can be printed side by side without translating field names. Tokens match exactly between sequential and supervisor here — expected, since this `if`-based router adds real wall-clock overhead (the extra node hops) but zero extra tokens, because it never calls the model to decide. That's the exact gap Approach 2 closes.

Approach 2 — an LLM-based supervisor, so the routing decision's real cost shows up

Story: an `if`-based router is free, so it hides the real cost of a supervisor — in production the supervisor *asks the model* where to go. This version pays for its decisions, so the numbers are honest. **If not:** the comparison would say "supervisors cost nothing extra", which is exactly wrong.

```
# architecture_comparison_practice.py – Real-world section
import time

def supervisor_llm(state):
    start = time.perf_counter()
    prompt = (
        f"Task: {state['task']}\n"
        f"Research done: {bool(state.get('research_text'))}\n"
        f"Write done: {bool(state.get('summary'))}\n"
        "Reply with exactly one word: RESEARCH, WRITE, or DONE."
    )
```

```

response = model.invoke(prompt)
elapsed = time.perf_counter() - start
decision = response.content.strip().upper()
# why: the routing call's own tokens and time count as real cost
call_tokens = response.usage_metadata["total_tokens"]

updates = {
    "total_tokens": state["total_tokens"] + call_tokens,
    "total_seconds": state["total_seconds"] + elapsed,
}

if decision == "RESEARCH":
    updates["routing_log"] = state["routing_log"] + ["research"]
    return Command(update=updates, goto="research_node")
if decision == "WRITE":
    updates["routing_log"] = state["routing_log"] + ["write"]
    return Command(update=updates, goto="write_node")
return Command(update=updates, goto=END)

builder_llm = StateGraph(SupervisorState)
builder_llm.add_node("supervisor", supervisor_llm)
builder_llm.add_node("research_node", research_node)
builder_llm.add_node("write_node", write_node)
builder_llm.add_edge(START, "supervisor")
graph_llm = builder_llm.compile()

def run_supervisor_llm(topic: str) -> dict:
    initial_state = {
        "task": topic, "research_text": "", "summary": "",
        "routing_log": [], "total_tokens": 0, "total_seconds": 0.0,
    }
    final_state = graph_llm.invoke(initial_state)
    return {
        "routing_log": final_state["routing_log"],
        "total_tokens": final_state["total_tokens"],
        "total_seconds": final_state["total_seconds"],
    }

```

```

# architecture_comparison_practice.py – Real-world section
if __name__ == "__main__":
    topics = ["climate change", "inflation", "sleep quality"]
    for topic in topics:
        seq = run_sequential(topic)
        sup_if = run_supervisor(topic)

```

```

sup_llm = run_supervisor_llm(topic)

seq_t, seq_s = seq["total_tokens"], seq["total_seconds"]
if_t, if_s = sup_if["total_tokens"], sup_if["total_seconds"]
llm_t, llm_s = sup_llm["total_tokens"], sup_llm["total_seconds"]

print(f"\n{topic}")
print(f"  sequential:      {seq_t}t {seq_s:.2f}s")
print(f"  supervisor (if):   {if_t}t {if_s:.2f}s")
print(f"  supervisor (llm):   {llm_t}t {llm_s:.2f}s")
print(f"    routed: {sup_llm['routing_log']}")

```

Expected output (exact numbers vary by run):

```

climate change
sequential:      306t 2.75s
supervisor (if):  306t 2.81s
supervisor (llm): 341t 3.62s
  routed: ['research', 'write']

```

The LLM-based supervisor calls the model 3 times total (route, route, route-to-done) — roughly 35 extra tokens and 0.8 extra seconds here, entirely from routing decisions that produce zero user-visible work. That's the real, honest cost `paper_design`'s prediction was estimating — an `if`-based router hides it completely.

Approach 3 — a written comparison table across all 3 designs

Story: three designs, three topics, one table — the form a skeptical reviewer can check at a glance.

If not: the result would be scattered across printouts, and the percentage you defend in Project 4 would be hand-counted.

```

# architecture_comparison_practice.py – Real-world section
def compare_all(topics: list[str]) -> None:
    header = f"{'topic':<16} {'sequential':>18} {'if':>18} {'llm':>18}"
    print(header)
    for topic in topics:
        seq = run_sequential(topic)
        sup_if = run_supervisor(topic)
        sup_llm = run_supervisor_llm(topic)

        seq_cell = f"{seq['total_tokens']}t/{seq['total_seconds']:.1f}s"
        if_cell = f"{sup_if['total_tokens']}t/{sup_if['total_seconds']:.1f}s"
        llm_cell = f"{sup_llm['total_tokens']}t/{sup_llm['total_seconds']:.1f}s"

        row = f"{topic:<16} {seq_cell:>18} {if_cell:>18} {llm_cell:>18}"

```

```
print(row)

if __name__ == "__main__":
    compare_all(["climate change", "inflation", "sleep quality"])
```

Expected output (exact numbers vary by run):

topic	sequential	if	llm
climate change	306t/2.8s	306t/2.8s	341t/3.6s
inflation	286t/2.6s	286t/2.7s	319t/3.4s
sleep quality	325t/2.9s	325t/3.0s	362t/3.8s

Real result, versus the `paper_design` guess: sequential and the `if`-based supervisor are nearly identical on tokens (routing there was free) and very close on speed (a few extra node hops, no extra model calls). The LLM-based supervisor is the one that actually matches `paper_design`'s prediction — roughly 10-15% more tokens and 25-30% more wall-clock time than sequential, entirely from a routing decision that, for this fixed two-step task, never had a real choice to make. `paper_design`'s verdict holds: for a task whose order never changes, a supervisor buys nothing except its own overhead.

Difference between Approach 1 and Approaches 2/3: Approach 1's comparison only shows sequential against a free-routing supervisor, which understates the pattern's real cost — a production supervisor calls the model to route, it doesn't use `if` checks. Approach 2 adds that missing, honest routing cost by building a second supervisor variant that actually pays for its decision. Approach 3 doesn't add new measurement — it lines up all 3 designs' numbers in one table across multiple topics, which is what turns "the LLM supervisor seemed to cost more" into a specific, repeatable percentage you could put in front of the tough senior engineer this document's "How to Practice" step asks you to defend against.

Which one should you actually write? For any real routing decision that could genuinely go more than one way, Approach 2's LLM-based supervisor is the one whose numbers you can trust — an `if`-based router only tells the truth about tasks so simple they didn't need a supervisor in the first place. For a task this fixed, though, the honest conclusion isn't "use the LLM supervisor" — it's "don't use a supervisor at all here," which is exactly what `paper_design` predicted before any code existed. Keep both supervisor variants in your toolkit: `if`-based routing for genuinely fixed pipelines (call it "sequential with extra steps," not really a supervisor), and LLM-based routing only once the next specialist truly can't be known ahead of time — which is what `ambiguous_routing` builds next.