

LangGraph Debugging — Worked Example: Solution

[Back to the exercise](#) · [Round 1: Basic](#) · [Round 2: Intermediate](#) · [Round 3: Real-world](#) · [Round 4: Multi-agent](#) · [Hints](#)

Story — `langgraph_debugging_practice.py` : graph bugs live in the wiring — route names, how state is merged, when a loop stops — not in any one node's code. Each fix here comes with a test that runs a tiny real graph, or checks the routing function directly, with stand-in nodes and no model. **If not:** you'd only see these bugs in a full, slow, paid run, and blame the model for a wiring mistake.

Every fix and test below goes in `practice/langgraph_debugging_practice.py`, and runs with `pytest langgraph_debugging_practice.py -v` from inside `practice/`. None of the tests need an API key.

- [Round 1: Basic](#)
- [Round 2: Intermediate](#)
- [Round 3: Real-world](#)
- [Round 4: Multi-agent](#)
- [Real cause vs. symptom fix](#)

[Back to the exercise](#) · [Round 1: Basic](#) · [Round 2: Intermediate](#) · [Round 3: Real-world](#) · [Round 4: Multi-agent](#) · [Hints](#)

Round: Basic

Real cause: `route()` returns `"finish"`, but the mapping given to `add_conditional_edges()` only has the keys `"call_tool"` and `"done"`. LangGraph looks the returned word up in the mapping, doesn't find it, and raises `KeyError: 'finish'` — a plain naming mismatch between the routing function and its mapping. This is Doc09's own Edge cases exercise (`unhandled_routing_value`).

Story: a bare `KeyError: 'finish'` gives almost nothing to go on, so this round trains searching your code for that exact word to find who produces it and who's supposed to accept it. **If not:** you'd go looking inside LangGraph for a bug that's a one-word mismatch in your own file.

The fix:

```
from typing import TypedDict
from langgraph.graph import StateGraph, START, END

class AgentState(TypedDict):
    needs_tool: bool

def agent_node(state):
    return {} # stand-in for the real agent

def tool_node(state):
    return {"needs_tool": False} # stand-in: the tool ran, done now

def route(state):
    if state["needs_tool"]:
        return "call_tool"
    # why: must be a key in the mapping below – "done", not "finish"
    return "done"

def build_graph():
    graph = StateGraph(AgentState)
    graph.add_node("agent", agent_node)
```

```

graph.add_node("tool_node", tool_node)
graph.add_edge(START, "agent")
graph.add_conditional_edges("agent", route, {
    "call_tool": "tool_node",
    "done": END,
})
graph.add_edge("tool_node", "agent")
return graph.compile()

```

Test that would have caught it:

```

def test_graph_can_take_the_done_branch():
    result = build_graph().invoke({"needs_tool": False})
    assert result["needs_tool"] is False

def test_graph_can_take_the_tool_branch_and_finish():
    result = build_graph().invoke({"needs_tool": True})
    assert result["needs_tool"] is False

```

One test per branch: a suite that only ever tested the "keep going" branch would never reach the one with the typo.

[Back to the exercise](#) · [Round 1: Basic](#) · [Round 2: Intermediate](#) · [Round 3: Real-world](#) · [Round 4: Multi-agent](#) · [Hints](#)

Round: Intermediate

Real cause: for a plain state field (no reducer), each node's update simply **replaces** the old value. `calculate_node`'s `{"findings": {"calc": 42}}` doesn't merge with `search_node`'s earlier `{"findings": {"search": ...}}` — it overwrites the whole dict, because nothing told LangGraph how to combine two updates to the same field.

Story: the log says "the model forgot" — but the state itself lost the data before the model ever saw it. This round trains checking state at each step before blaming reasoning. **If not:** you'd add "remember earlier findings" to the prompt, and nothing would change.

The fix:

```
from typing import Annotated, TypedDict
from langgraph.graph import StateGraph, START, END

def merge_findings(current, update):
    # why: a reducer tells LangGraph HOW to combine old and new —
    # here, keep every key from both dicts
    merged = dict(current)
    for key in update:
        merged[key] = update[key]
    return merged

class FindingsState(TypedDict):
    task: str
    findings: Annotated[dict, merge_findings]

def search_node(state):
    # stand-in for run_search(state["task"])
    return {"findings": {"search": {"refund_policy": "30 days"}}}

def calculate_node(state):
    # stand-in for run_calculation(state["task"])
    return {"findings": {"calc": 42}}
```

```
def build_findings_graph():
    graph = StateGraph(FindingsState)
    graph.add_node("search", search_node)
    graph.add_node("calculate", calculate_node)
    graph.add_edge(START, "search")
    graph.add_edge("search", "calculate")
    graph.add_edge("calculate", END)
    return graph.compile()
```

Test that would have caught it:

```
def test_findings_from_both_nodes_are_kept():
    result = build_findings_graph().invoke({"task": "x", "findings": {}})
    assert "search" in result["findings"]
    assert "calc" in result["findings"]
```

The test has to run both nodes through a real graph — only the graph applies the reducer. Calling the two node functions by hand and merging with `dict.update()` would skip the reducer and prove nothing.

[Back to the exercise](#) · [Round 1: Basic](#) · [Round 2: Intermediate](#) · [Round 3: Real-world](#) · [Round 4: Multi-agent](#) · [Hints](#)

Round: Real-world

Real cause: `route()` only checks confidence against a fixed threshold — there's no counter of how many search rounds have happened. Most questions pass 0.5 within a couple of rounds. A certain kind of ambiguous question stays "somewhat relevant, never enough" forever, and with nothing else to stop it, the loop runs to LangGraph's hard recursion limit — the *symptom* (a recursion error) fires 25 steps after the *real cause* (a loop with only one way out) let it start.

Story: `GraphRecursionError` has no frame from your own node, which is the clue: nothing crashed, the *routing* never let go. This round trains giving every loop two exits — "good enough" and "tried enough". **If not:** you'd raise the recursion limit, and the same questions would just loop longer and cost more.

The fix:

```
from typing import TypedDict

class SearchState(TypedDict):
    task: str
    confidence: float
    search_rounds: int

def route_search(state):
    if state["confidence"] >= 0.5:
        return "done"
    # why: the second way out – stop after 4 tries, confident or not
    if state["search_rounds"] >= 4:
        return "give_up"
    return "search_again"
```

Each search node adds 1 to `search_rounds` in its update, and `"give_up"` routes to a node that returns a clear "couldn't find a confident answer" result.

Test that would have caught it:

```
def test_search_loop_gives_up_after_four_rounds():
    state = {"task": "x", "confidence": 0.3, "search_rounds": 4}
```

```
assert route_search(state) == "give_up"
```

```
def test_search_loop_keeps_going_before_the_limit():  
    state = {"task": "x", "confidence": 0.3, "search_rounds": 1}  
    assert route_search(state) == "search_again"
```

The routing function is plain Python, so the "never confident" case can be tested with a hand-made state in milliseconds — no need to run 25 real search rounds.

[Back to the exercise](#) · [Round 1: Basic](#) · [Round 2: Intermediate](#) · [Round 3: Real-world](#) · [Round 4: Multi-agent](#) · [Hints](#)

Round: Multi-agent

Real cause: two bugs stacked on top of each other. First, `review_route` has the same problem as Round 3 — no cap on revision rounds. Second, the Reviewer keeps rejecting because the "compliance detail" it wants was never written into `research_findings` at all (a Research-agent gap), so the Writer is asked, every round, for something it can't produce. Neither shows up testing the Writer/Reviewer pair with hand-written examples, because those always include findings the Writer *can* satisfy.

Story: the loop has two causes, and each fix covers a different one — the cap stops the endless loop no matter why, and looking one layer back finds why it looped. This round trains fixing both, not just the one you saw first. **If not:** the cap alone would turn "loops forever" into "gives up every time", and the missing compliance detail would stay missing.

The fix:

```
from typing import TypedDict

class ReviewState(TypedDict):
    task: str
    research_findings: str
    draft: str
    review_passed: bool
    revision_count: int

def count_revision(state):
    # how: the Writer node adds this to its update on every draft
    return {"revision_count": state["revision_count"] + 1}

def review_route(state):
    if state["review_passed"]:
        return "done"
    # why: Doc11's Build Task rule – a hard limit on revision rounds
    if state["revision_count"] >= 3:
        return "give_up"
    return "revise"
```

The graph maps `"give_up"` to a `give_up_report` node that says clearly the agents couldn't agree. Then go one layer back: make the Reviewer's rejection name what's missing from `research_findings`, so the gap in the Research agent's output is visible instead of hidden inside an endless loop.

Test that would have caught it:

```
def test_review_loop_gives_up_after_three_revisions():
    state = {"task": "x", "research_findings": "", "draft": "d",
            "review_passed": False, "revision_count": 3}
    assert review_route(state) == "give_up"

def test_each_draft_counts_as_one_revision():
    state = {"task": "x", "research_findings": "", "draft": "d",
            "review_passed": False, "revision_count": 1}
    assert count_revision(state) == {"revision_count": 2}
```

The first test proves the loop can end; the second proves the counter really moves, so the limit is actually reached.

[Back to the exercise](#) · [Round 1: Basic](#) · [Round 2: Intermediate](#) · [Round 3: Real-world](#) · [Round 4: Multi-agent](#) · [Hints](#)

Real cause vs. symptom fix

A symptom fix here would be wrapping `invoke()` in a `try/except KeyError` (Round 1), telling the model in the prompt to "remember all earlier findings" (Round 2), or raising `recursion_limit` to 100 (Rounds 3 and 4). Each makes one visible failure go away while the wiring stays wrong: the mapping still doesn't match, state still overwrites itself, and loops still have only one exit — so the next unusual input hits the same wall, later and more expensively. The real fixes all change the graph's structure — matching route names, a reducer that says how to merge, and a second exit on every loop — which is why they hold for inputs you haven't tried yet.