

Multi-agent Debugging — Worked Example

[Back to the exercise](#) · [Round 1: Basic](#) · [Round 2: Intermediate](#) · [Round 3: Real-world](#) · [Round 4: Multi-agent](#) · [Solution](#)

One scenario, followed across all 4 rounds: Doc11's Project 4 supervisor (`agents/supervisor.py`), routing to Research → Analysis → Writer → Reviewer. For this category, "Round 4: Multi-agent" means the failure is only visible running the *entire* 4-agent pipeline together — not testing the supervisor plus one specialist in isolation. Work through the rounds in order. Each ends with "**What do you think is wrong?**" — stop and actually answer before reading on.

- [Round 1: Basic](#)
- [Round 2: Intermediate](#)
- [Round 3: Real-world](#)
- [Round 4: Multi-agent](#)

[Back to the exercise](#) · [Round 1: Basic](#) · [Round 2: Intermediate](#) · [Round 3: Real-world](#) · [Round 4: Multi-agent](#) · [Solution](#)

Round: Basic

Setup: the supervisor node hands off to a specialist using `Command` :

```
def supervisor_node(state):
    next_agent = decide_next_agent(state)
    return Command(goto=next_agent)
```

`decide_next_agent()` returns one of a fixed set of strings, one of which has a typo baked in:

```
def decide_next_agent(state):
    ...
    return "Reserach"    # the typo
```

Symptoms: no exception at all. Every run that should go to the research specialist finishes in about a second — with no research, no draft, and an empty result.

Log output:

```
[supervisor] routing to: Reserach
Task supervisor with path ('__pregel_pull', 'supervisor') wrote to
unknown channel branch:to:Reserach, ignoring it.
[result] research_findings = ""    draft = ""
```

Expected vs. actual:

- Expected: `Command(goto="research_agent")` hands control to the Research agent.
- Actual: LangGraph logs one warning line and ends the run quietly — nothing raises, so nothing in your `try/except` or your error alerts ever sees it.

What do you think is wrong?

[Back to the exercise](#) · [Round 1: Basic](#) · [Round 2: Intermediate](#) · [Round 3: Real-world](#) · [Round 4: Multi-agent](#) · [Solution](#)

Round: Intermediate

Setup: routing is fixed — the Research agent runs correctly and writes its output:

```
def research_node(state):
    finding = run_research_tools(state["task"])
    return Command(goto="analysis_agent", update={"finding": finding})
```

The shared state schema, defined separately, actually declares a different field name:

```
class SharedState(TypedDict):
    task: str
    research_findings: str
    analysis: str
    draft: str
```

Symptoms: no exception anywhere. The graph runs start to finish and produces a final answer — just a noticeably weak, generic one, every time.

Log output:

```
[research_agent] finding = "The refund policy allows returns within 30 days
of purchase."
[analysis_agent] state.research_findings = ""
[analysis_agent] WARNING: no research findings to analyze,
proceeding with general knowledge
[writer_agent] draft: "Our return policy details may vary;
please check with support."
```

Expected vs. actual:

- Expected: the Analysis agent reads what Research actually found and builds on it.
- Actual: the Analysis agent's view of `research_findings` is always an empty string, even though Research clearly found something real, one line above in the same log.

What do you think is wrong?

[Back to the exercise](#) · [Round 1: Basic](#) · [Round 2: Intermediate](#) · [Round 3: Real-world](#) · [Round 4: Multi-agent](#) · [Solution](#)

Round: Real-world

Setup: the field-name bug is fixed. The supervisor picks between two specialists using an LLM call, given each specialist's description:

```
SPECIALISTS = {  
    "research_agent": "Finds factual information and data to answer a "  
                      "question.",  
    "analysis_agent": "Analyzes information and data to draw conclusions.",  
}
```

Symptoms: only sometimes happens — most requests route to the obviously-correct specialist. A specific style of request routes inconsistently: the same *kind* of task, worded slightly differently, sometimes goes to Research and sometimes to Analysis.

Log output (5 similar requests, same underlying task):

```
[supervisor] "find and interpret last quarter's revenue numbers"  
-> analysis_agent  
[supervisor] "look up and interpret last quarter's revenue numbers"  
-> research_agent  
[supervisor] "get and interpret last quarter's revenue figures"  
-> analysis_agent  
[supervisor] "find last quarter's revenue numbers and interpret them"  
-> research_agent  
[supervisor] "interpret last quarter's revenue numbers after finding them"  
-> analysis_agent
```

Expected vs. actual:

- Expected: a supervisor routing decision is a deliberate design choice with a knowable reason, not a coin flip.
- Actual: for tasks that genuinely need both "finding" and "interpreting," the supervisor's routing flips unpredictably between the two specialists depending on exact wording — because both specialist descriptions genuinely apply, and the supervisor has no tiebreaker.

What do you think is wrong?

[Back to the exercise](#) · [Round 1: Basic](#) · [Round 2: Intermediate](#) · [Round 3: Real-world](#) · [Round 4: Multi-agent](#) · [Solution](#)

Round: Multi-agent

This round only shows up running the full 4-agent pipeline — not the supervisor alone, not any 2-agent pair alone.

Setup: the Round 2 field-name bug got "fixed" quickly under deadline pressure by adding a second field instead of renaming the first one:

```
class SharedState(TypedDict):
    task: str
    research_findings: str    # never actually written to, still
    finding: str              # what research_node actually writes to
    analysis: str
    draft: str
    review_passed: bool
    revision_count: int
```

`analysis_node` was updated to read from `finding` and works correctly now. `reviewer_node`, written independently by someone checking the draft against the *original* field name, wasn't:

```
def reviewer_node(state):
    findings = state["research_findings"]
    if findings == "" or findings not in state["draft"]:
        return Command(goto="writer_agent", update={
            "review_passed": False,
            "revision_count": state["revision_count"] + 1,
        })
    return Command(goto=END, update={"review_passed": True})
```

Symptoms: testing Research + Analysis together: works perfectly, `finding` flows through correctly. Testing Writer + Reviewer together with a hand-written fake `draft` and a hand-written fake `research_findings` that match: works perfectly. Running all 4 agents together on a real task: the Reviewer rejects every draft, forever, until it hits the revision cap and gives up.

Log output:

```
[research_agent] finding = "Q1 revenue grew 12% year over year."
[analysis_agent] analysis = "Growth is driven mainly by the new product line."
```

```
[writer_agent] draft: "Q1 revenue grew 12% year over year, driven mainly
    by the new product line."
[reviewer] state.research_findings = ""    (still the unused field)
[reviewer] REJECTED: draft does not reflect research findings
[writer_agent] revision #2: draft: "Q1 revenue grew 12% year over year,
    driven mainly by the new product line."
[reviewer] REJECTED: draft does not reflect research findings
... (continues to revision_count = 3, then gives up) ...
```

Expected vs. actual:

- Expected: the Reviewer checks the draft against the same findings the Writer actually used, and approves a genuinely accurate draft.
- Actual: Research and Analysis were fixed to use `finding`, but the Reviewer still reads the old `research_findings` field — which nobody writes, so it's always empty — and its check treats empty findings as "the draft doesn't reflect the findings". So it rejects every draft. Neither pair test could catch this, because each one used whichever field name its own author picked.

What do you think is wrong?

[Back to the exercise](#) · [Round 1: Basic](#) · [Round 2: Intermediate](#) · [Round 3: Real-world](#) · [Round 4: Multi-agent](#) · [Solution](#)

Full solution: [Show me the solution](#)