

Basic (critique your own Stage-1 design) — Hints

[Back to the exercise](#) · [Hint 1](#) · [Hint 2](#) · [Solution](#)

Only 2 hints — work through them in order against your own written design before reading ahead. Each hint has 3 depth levels: **Basic** (the plain framing — what to ask), **Intermediate** (the real technique, in proper design-review vocabulary), **Advanced** (the senior-level judgment call — the trap in doing this review on your own work). Read Basic first even if you already know the technique — it's the fastest way to see exactly what each deeper level adds.

- [Hint 1 — The idea, and the exact check](#)
- [Hint 2 — The plan, and a worked example](#)

[Back to the exercise](#) · [Hint 1](#) · [Hint 2](#) · [Solution](#)

Hint 1 — The idea, and the exact check

Basic Version

You already wrote a design before this exercise — a list of agents, tools, and pieces you plan to build. This exercise asks you to go back through that list and, for each piece, ask one plain question: "what would break if I removed this?"

If nothing you can name would actually break, that piece is probably there because it seemed like a good idea, not because the requirements needed it. That's the exact thing this exercise is meant to catch, before you spend time building it.

[Back to the exercise](#) · [Hint 1](#) · [Hint 2](#) · [Solution](#)

Intermediate Version

Doc16 frames this directly: "requirements → architecture is matching, not inspiration." The self-review technique is to build a **traceability table** — walk your design's component list and, for each one, force yourself to name the *specific* requirement (not a general good practice) it satisfies. Two failure patterns to watch for while doing this:

- **Overbuilt complexity** — a piece justified by "this is what real production systems have" rather than by anything in your actual requirements. A caching layer for traffic you don't have yet is the canonical example.
- **Solving a problem you don't have** — a piece defends against a failure mode your requirements never mentioned as a real risk.

Then run the check in the reverse direction too: for each *non-functional* requirement (speed, cost, reliability target), confirm at least one component actually addresses it. A requirement with nothing satisfying it is a gap, not just noise — this catches under-building, which is just as common as over-building and easier to miss because nothing "looks wrong" about a design that's simply incomplete.

Go through your own component list, both directions, before Hint 2.

[Back to the exercise](#) · [Hint 1](#) · [Hint 2](#) · [Solution](#)

Advanced Version

Here's the trap in doing this review on your own design: you're the same person who chose every component, so you're also the most motivated person alive to find a requirement that justifies keeping it. A traceability table filled out by the person who built the thing it's tracing isn't automatically honest — it can be gamed without you noticing, by quietly writing the "requirement" column to match the component you already wanted, instead of checking the component against a requirement that existed independently first.

The real advanced self-review question isn't just "can I name a requirement for this" — it's **"did I write this requirement down before or after I built the component it's justifying?"** If you can't remember requiring redaction before you designed a redaction layer, and you're now retroactively pointing at "must never expose raw customer data" to justify it, that's not a failure of the check — that's the check working, catching a rationalization instead of a match. The honest fix isn't to delete the requirement (it's probably real) — it's to be suspicious of any row you filled in *fast* and *confidently*, since post-hoc justifications come easily precisely because you're good at constructing them.

A concrete way to catch this: for your 2-3 most confident rows, try to state the requirement in a form you'd have written *before* seeing your own design — if you can't, or the wording suspiciously echoes your component's own name, treat that row as unverified until you check it against the actual written requirements document, not your memory of it.

Difference between Basic, Intermediate, and Advanced: Basic asks "what breaks if I remove it," which catches the obviously unjustified pieces. Intermediate turns that into a real two-direction traceability check, catching both over-building and under-building. Advanced questions the check itself — a self-review has a built-in bias problem the two-direction table doesn't fix on its own, and catching your own rationalizations is the harder, more senior half of this exercise.

[Back to the exercise](#) · [Hint 1](#) · [Hint 2](#) · [Solution](#)

Hint 2 — The plan, and a worked example

Basic Version

```
list every piece in your design (agents, tools, database, etc.)

for each piece, write down:
    which requirement made me add this?

if you can't answer that:
    is it "just in case" (overbuilt) or "fixing something that isn't
    actually a risk here" (solving a problem you don't have)?
    decide: cut it, or write down why you're keeping it anyway
```

[Back to the exercise](#) · [Hint 1](#) · [Hint 2](#) · [Solution](#)

Intermediate Version

build a traceability table: Component | Requirement it satisfies

for each component:

- write the specific requirement (not a general principle)

- if empty/vague: classify as overbuilt complexity, or a fix for a non-existent risk

- decide: cut, defer with a written note, or find the real requirement you missed

check the reverse direction too:

- for each NON-FUNCTIONAL requirement (speed, cost, reliability target), confirm at least one component in your design actually addresses it

- a requirement with nothing satisfying it is a gap

Worked example row:

Component: a caching layer in front of the retriever

Stated requirement: "must handle 50,000 requests/day, p95 latency under 2 seconds"

Verdict: KEEP -- the latency target is real and stated, and a cache directly serves it.

Component: a second, redundant fact-checker agent running in parallel with the first

Stated requirement: none found

Verdict: CUT (or defer with a note) -- "more checking seems safer" is not the same as a requirement calling for it.

[Back to the exercise](#) · [Hint 1](#) · [Hint 2](#) · [Solution](#)

Advanced Version

The same worked example, with the bias check added on top:

Component: field-level redaction before any data reaches the LLM
 Stated requirement (as I'd write it now): "must never expose raw customer data to the LLM"

Bias check: did I write this requirement down BEFORE designing the redaction layer, or am I pointing at it now because it happens to match what I already built?

-> if this was in the live requirements handoff, it's real: KEEP, confidently.

-> if I can't recall it being stated, and I'm inferring it from "well, customer data is sensitive, so obviously..." -- that's MY reasoning, not a requirement. Flag it: still probably the right call, but check it against the actual requirements document before treating it as settled.

Component: an admin dashboard with 12 chart types

Stated requirement: "watching/monitoring" (generic, from the build list)

Bias check: this one is easy to catch -- "watching/monitoring" doesn't specify 12 chart types, or any number. The gap between the vague ask and the specific thing built is the tell.

Verdict: CUT down to 3-4 charts tied to actual non-functional targets.

Run the bias check against your own most confident row first — that's usually where a rationalization hides, precisely because it felt too obvious to question. Then check the [Solution](#).

Difference between Basic, Intermediate, and Advanced: Basic and Intermediate both trust the table once it's filled in. Advanced adds a pass that questions the table itself — checking whether each "requirement" was real and independent, or quietly reverse-engineered to fit a component you'd already decided to keep.

[Back to the exercise](#) · [Hint 1](#) · [Hint 2](#) · [Solution](#)

Full solution: [Show me the solution](#)