Step 2 — A Real ReAct Agent That Calls a Discovered Tool — Solution

> Back to the step · Hint 1 · Hint 2

Story: the point of the whole project — an agent that acts on tools it never had written into it. The bridge turns discovered MCP tools into OpenAI's function shape; the loop sends them to the model and runs whatever the model picks, through the MCP client. **If not:** you'd be back to hardcoding a `@tool` function for every capability — exactly what MCP exists to avoid.

`src/tool_bridge.py`

Story: two small adapters. `mcp_tools_to_openai_schema` turns each discovered tool into the dict OpenAI's `tools` parameter wants — an MCP tool's `input_schema` is already JSON Schema, so it's mostly a copy. `extract_text` pulls the text out of a `CallToolResult`, which is a list of content blocks, not a plain string. **If not:** the model call would reject the tool shape, and printing a result would crash on an object that isn't a string.

```
# Adapters between MCP's shapes and OpenAI's. Kept in one file so the loop
# never has to know either library's exact types.
from src.utils.logger import get_logger


def one_openai_tool(namespaced_name, description, input_schema):
    return {
        "type": "function",
        "function": {
            "name": namespaced_name,
            "description": description,
            "parameters": input_schema,
        },
    }


def mcp_tools_to_openai_schema(named_tools):
    logger = get_logger(__name__)
    schemas = []
    for namespaced_name, tool in named_tools:
        description = tool.description
```

```
        if description is None:
            description = ""
        schemas.append(
            one_openai_tool(namespaced_name, description, tool.input_schema)
        )
    logger.info(f"converted {len(schemas)} MCP tools to OpenAI schema")
    return schemas


def extract_text(call_result):
    parts = []
    for block in call_result.content:
        text = getattr(block, "text", None)
        if text is not None:
            parts.append(text)
    return "\n".join(parts)
```

`named_tools` is the list of `(namespaced_name, mcp_tool)` pairs the pool already built in Step 1 — the bridge never touches a raw server, only what was discovered.

`src/server_pool.py` — two methods added

Story: the pool now also hands out the OpenAI-shaped tool list and dispatches a call by its namespaced name. The loop calls these two methods and nothing else. **If not:** the agent would need to know about sessions and routing, instead of just "here are the tools, call this one."

Add these two methods to the `MCPServerPool` class from Step 1:

```
    def openai_tools(self):
        return mcp_tools_to_openai_schema(self.named_tools)

    async def call_tool(self, namespaced_name, arguments):
        server_name, tool_name = self.routing[namespaced_name]
        session = self.sessions[server_name]
        result = await session.call_tool(tool_name, arguments)
        return extract_text(result)
```

and the imports at the top of the file:

```
from src.tool_bridge import extract_text, mcp_tools_to_openai_schema
```

(Step 4 wraps `call_tool`'s dispatch in a `try/except`; for now a clean call is enough.)

`src/models/llm.py`

Story: the one file that talks to OpenAI — create the client, and one `ask_model()` that sends the conversation plus the discovered tools and returns the model's message. **If not:** OpenAI calls would be scattered through the loop, and swapping the model or adding a startup check would mean editing several files.

```
# The only file that talks to the OpenAI SDK: the client, a startup key
# check, and the one tool-calling request the ReAct loop makes.
import openai
from openai import OpenAI

MODEL = "gpt-4o-mini"

def create_client(config):
    return OpenAI(api_key=config.openai_api_key)

def check_api_key_at_startup(client):
    try:
        client.chat.completions.create(
            model=MODEL,
            messages=[{"role": "user", "content": "ping"}],
            max_tokens=1,
        )
    except openai.AuthenticationError:
        return False
    return True

def ask_model(client, messages, tools):
    response = client.chat.completions.create(
        model=MODEL, messages=messages, tools=tools
    )
    return response.choices[0].message
```

`src/agent.py`

Story: the hand-built ReAct loop (no `create_agent` / `AgentExecutor`) — ask the model, run any tool calls through the pool, append each result as a `tool` message, repeat until the model answers with no tool calls. A plain `max_iterations` cap stops a runaway loop. **If not:** a confused model could call tools forever, and the "discovered tools" idea would have no loop to prove it.

```
# The ReAct loop: think (ask the model) -> act (call a discovered tool via
# the pool) -> observe (feed the result back). Bounded by max_iterations.
```

```
import json
from dataclasses import dataclass

from src.exceptions import MaxIterationsExceeded
from src.models.llm import ask_model
from src.utils.logger import get_logger


@dataclass
class AgentResult:
    answer: str
    steps: list


def build_messages(task):
    system = (
        "You are a helpful agent. Use the available tools to answer the "
        "task. When you have enough information, answer in plain words."
    )
    return [
        {"role": "system", "content": system},
        {"role": "user", "content": task},
    ]


async def run_agent(task, pool, client, max_iterations=6):
    logger = get_logger(__name__)
    messages = build_messages(task)
    tools = pool.openai_tools()
    steps = []

    for step_number in range(max_iterations):
        reply = ask_model(client, messages, tools)
        if not reply.tool_calls:
            logger.info("agent finished with a final answer")
            return AgentResult(answer=reply.content, steps=steps)

        messages.append({
            "role": "assistant",
            "content": reply.content,
            "tool_calls": _tool_calls_as_dicts(reply.tool_calls),
        })
        for call in reply.tool_calls:
            name = call.function.name
            arguments = json.loads(call.function.arguments)
            logger.info(f"step {step_number}: calling {name} {arguments}")
            observation = await pool.call_tool(name, arguments)
```

```
            steps.append({"tool": name, "args": arguments,
                          "observation": observation})
            messages.append({
                "role": "tool",
                "tool_call_id": call.id,
                "content": observation,
            })

    raise MaxIterationsExceeded(
        f"agent did not finish within {max_iterations} steps"
    )


def _tool_calls_as_dicts(tool_calls):
    out = []
    for call in tool_calls:
        out.append({
            "id": call.id,
            "type": "function",
            "function": {
                "name": call.function.name,
                "arguments": call.function.arguments,
            },
        })
    return out
```

`main.py` — now runs the agent

```
# Step 2 demo: connect to one server, discover its tools, and run the
# agent on a real task — printing each tool call and the final answer.
import asyncio
import sys

from src.agent import run_agent
from src.config import load_config
from src.exceptions import MissingConfigError
from src.models.llm import create_client
from src.server_pool import MCPServerPool
from src.servers_config import SERVERS

TASK = "What does my note about the Q3 roadmap say?"


async def run_demo():
    try:
```

```
        config = load_config()
    except MissingConfigError as e:
        print(f"Setup problem: {e}")
        sys.exit(1)

    client = create_client(config)
    pool = MCPServerPool(SERVERS, connect_timeout=config.connect_timeout)
    await pool.connect_all()
    print(f"Connected servers: {list(pool.sessions.keys())}")
    print(f"Tools discovered: {len(pool.named_tools)}")
    try:
        print(f"\nTask: {TASK}")
        result = await run_agent(
            TASK, pool, client, max_iterations=config.max_iterations
        )
        for step in result.steps:
            print(f"  -> {step['tool']} {step['args']}")
        print(f"Answer: {result.answer}")
    finally:
        await pool.aclose()


def main():
    asyncio.run(run_demo())


if __name__ == "__main__":
    main()
```

Expected output (`python main.py` , with a real key) — the chained case: `search_notes` finds the note id, then `read_note` reads it using that id.

```
Connected servers: ['notes']
Tools discovered: 2

Task: What does my note about the Q3 roadmap say?
  -> notes__search_notes {'query': 'roadmap'}
  -> notes__read_note {'note_id': 'n1'}
Answer: Your Q3 roadmap note says to ship the billing API first, then the
mobile app beta.
```

`tests/unit/test_tool_bridge.py`

```
# Checks the schema bridge with tiny fakes — no server, no network.
# Run from the repo root: python -m tests.unit.test_tool_bridge
from src.tool_bridge import extract_text, mcp_tools_to_openai_schema


class FakeTool:
    def __init__(self, name, description, input_schema):
        self.name = name
        self.description = description
        self.input_schema = input_schema


class FakeBlock:
    def __init__(self, text):
        self.text = text


class FakeResult:
    def __init__(self, texts):
        self.content = []
        for text in texts:
            self.content.append(FakeBlock(text))


def test_schema_shape():
    tool = FakeTool("search_notes", "Search notes.",
                    {"type": "object", "properties": {}})
    schemas = mcp_tools_to_openai_schema([("notes__search_notes", tool)])
    one = schemas[0]
    result = "OK"
    if one["type"] != "function":
        result = "WRONG"
    if one["function"]["name"] != "notes__search_notes":
        result = "WRONG"
    if one["function"]["description"] != "Search notes.":
        result = "WRONG"
    if "properties" not in one["function"]["parameters"]:
        result = "WRONG"
    print(f"case 1 (openai schema shape)   -> {result}")


def test_missing_description():
    tool = FakeTool("t", None, {"type": "object"})
    schemas = mcp_tools_to_openai_schema([("s__t", tool)])
    result = "OK"
    if schemas[0]["function"]["description"] != "":
        result = "WRONG"
```

```
    print(f"case 2 (None description -> '') -> {result}")


def test_extract_text_joins():
    result_obj = FakeResult(["line one", "line two"])
    text = extract_text(result_obj)
    result = "OK"
    if text != "line one\nline two":
        result = "WRONG"
    print(f"case 3 (extract_text joins)    -> {result}")


if __name__ == "__main__":
    test_schema_shape()
    test_missing_description()
    test_extract_text_joins()
```

Expected output:

```
case 1 (openai schema shape)   -> OK
case 2 (None description -> '') -> OK
case 3 (extract_text joins)    -> OK
```

`tests/integration/test_agent.py` (cases 1–2)

Story: the real proof — launch the example server over a real MCP connection and confirm the agent discovers its tools and chains two calls. **If not:** a unit test with fakes could pass while the real protocol path was broken.

```
# Integration test: launches the real example servers over MCP. Each case
# prints OK or WRONG. Run from the repo root:
#   python -m tests.integration.test_agent
import asyncio

from src.agent import run_agent
from src.config import load_config
from src.models.llm import create_client
from src.server_pool import MCPServerPool

NOTES = {"name": "notes", "command": "python",
         "args": ["servers/notes_server.py"]}


async def connect(servers):
```

```
    config = load_config()
    client = create_client(config)
    pool = MCPServerPool(servers, connect_timeout=config.connect_timeout)
    await pool.connect_all()
    return config, client, pool


async def test_discovery_namespaced():
    config, client, pool = await connect([NOTES])
    names = []
    for namespaced, tool in pool.named_tools:
        names.append(namespaced)
    result = "OK"
    if "notes__search_notes" not in names:
        result = "WRONG"
    if "notes__read_note" not in names:
        result = "WRONG"
    await pool.aclose()
    print(f"case 1 (tools discovered + namespaced) -> {result}")


async def test_chained_single_server():
    config, client, pool = await connect([NOTES])
    out = await run_agent("What does my note about the Q3 roadmap say?",
                          pool, client, config.max_iterations)
    result = "OK"
    if len(out.steps) != 2:
        result = "WRONG"
    if out.steps[0]["tool"] != "notes__search_notes":
        result = "WRONG"
    if out.steps[1]["tool"] != "notes__read_note":
        result = "WRONG"
    await pool.aclose()
    print(f"case 2 (chained search then read)      -> {result}")


async def main():
    await test_discovery_namespaced()
    await test_chained_single_server()


if __name__ == "__main__":
    asyncio.run(main())
```

Expected output (cases 3–4 are added in Steps 3 and 4):

```
case 1 (tools discovered + namespaced) -> OK
case 2 (chained search then read)      -> OK
```